# So you want to do some logging. . . (PT. 6 IIS Logs)

blog.iso365down.com/so-you-want-to-do-some-logging-pt-6-iis-logs-52f70819567a

HanSolo71 December 26, 2023

## Creating a New Index

Before moving forward we will want to create a new index in Graylog. Logs from web servers can be generate a lot of traffic and you may want to change retention based on storage needs.

To make a new index set, use the top menu and select **System > Indices.**

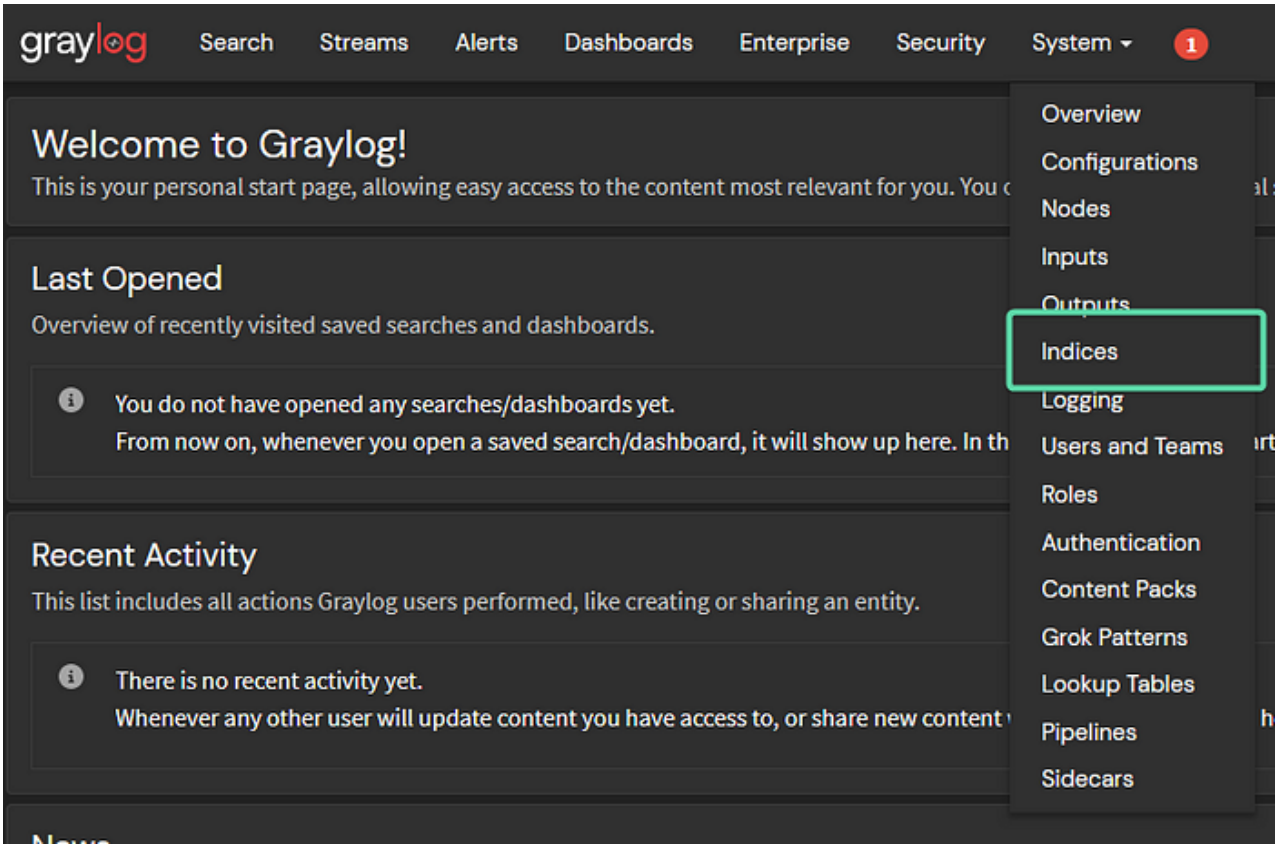


Time for a new index!

Then select **Create index set.**

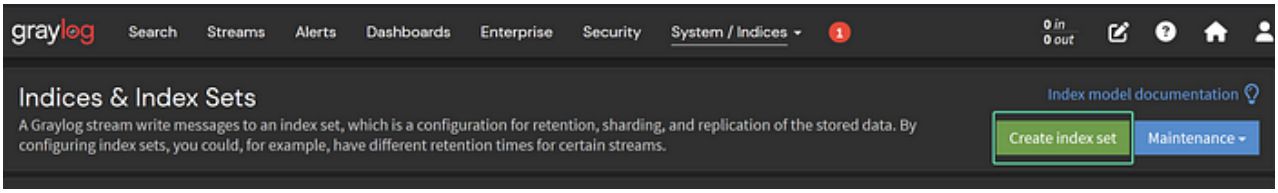


Make the new index

We will only need to worry about *Title* and *Index prefix* fields along with rotation strategy.

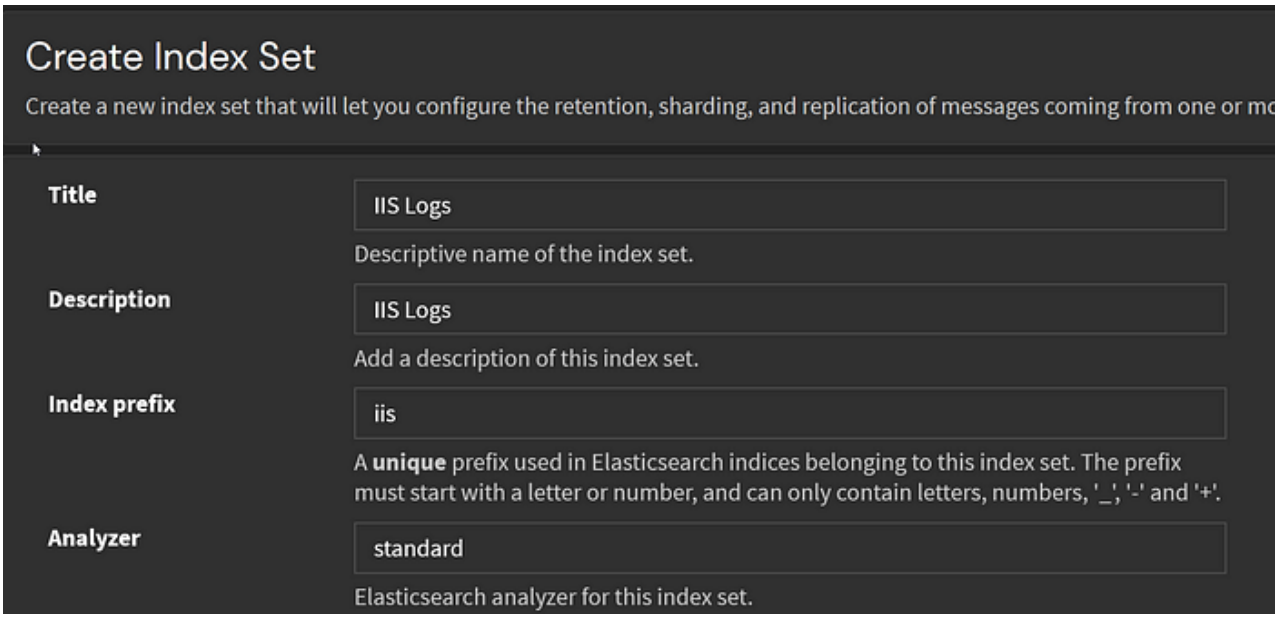


Name the new index set

We want to keep our IIS/Web access logs for 30 days.

To do this, we will use a rotation strategy of **Index Time** with a duration of **P1D**. We want to set the retention strategy to **Delete Index** and set the max number of indices to 30.

These setting will create a new index every day and delete the oldest index when index 31 is created

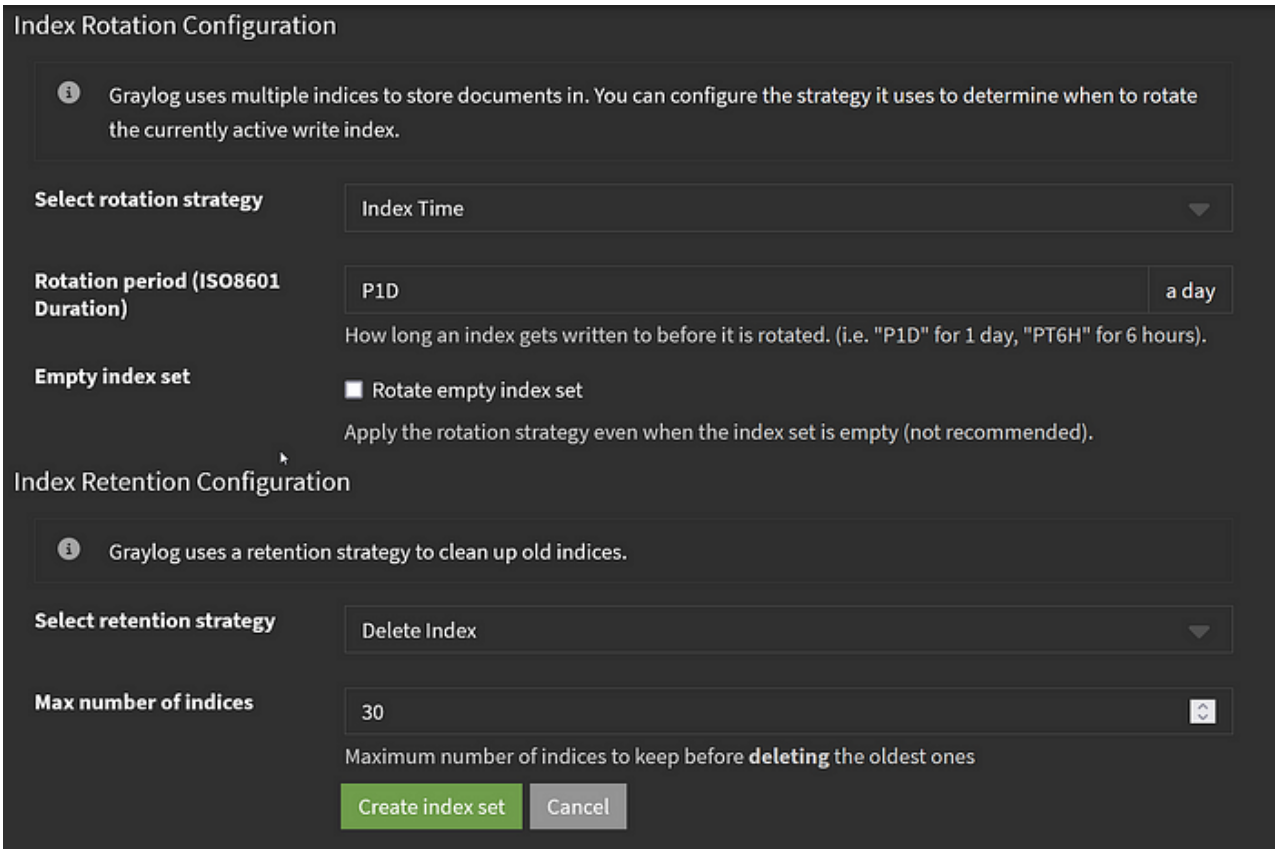


These setting will create a new index every day and delete the oldest index when index 31 is created

## Setting up IIS

Before moving forward we will want to enable all logging fields and ensure we are using W3C logs. We will start by opening the **IIS Manager.**

Select the IIS Servers Logging Rules

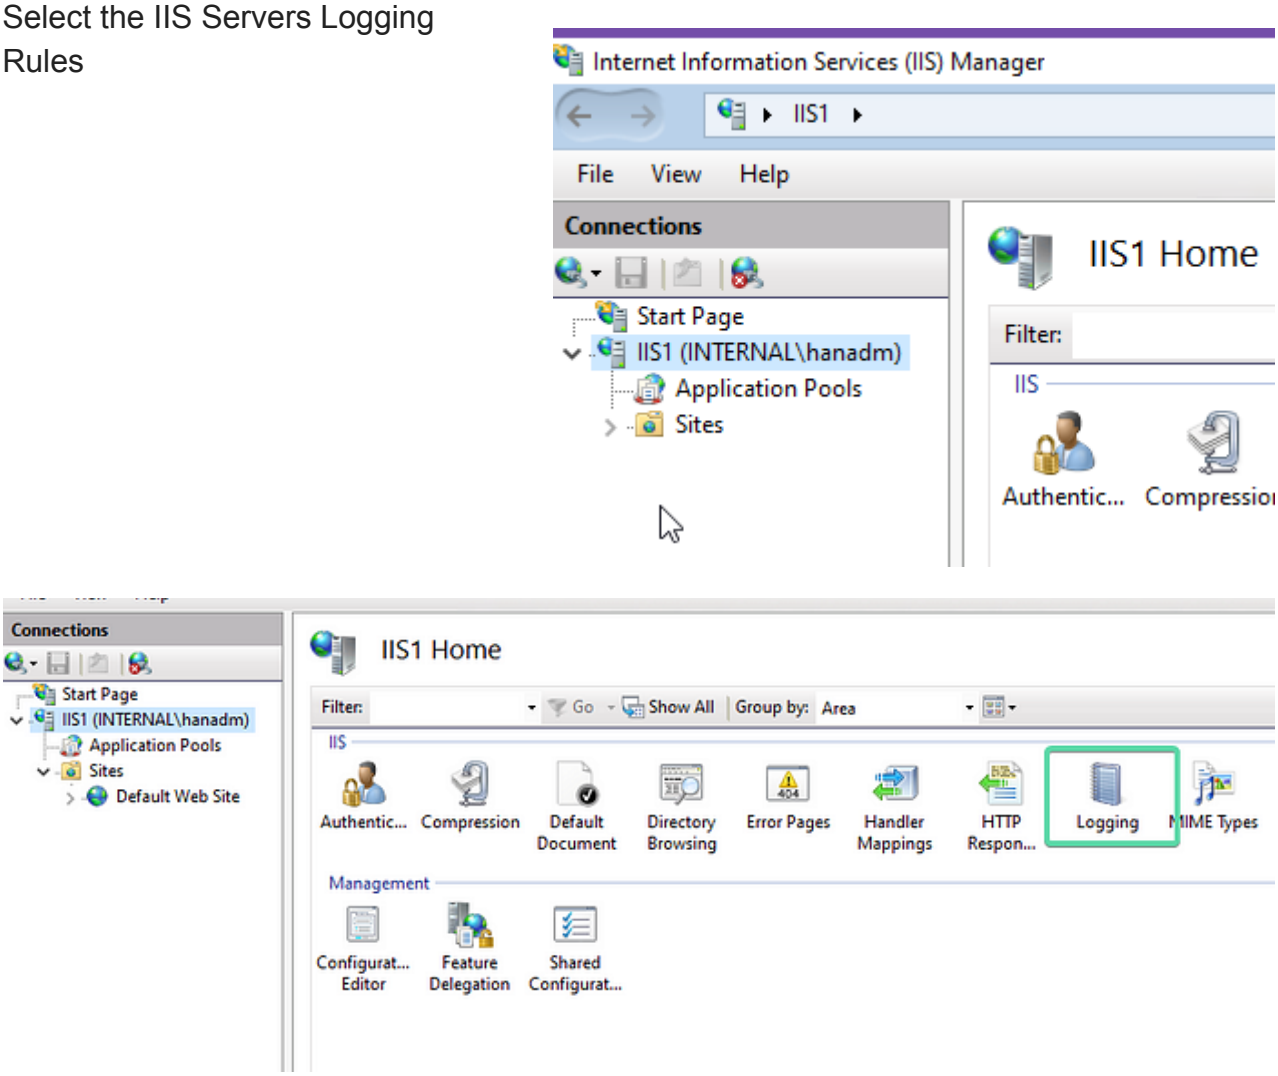


Edit logging

Ensure you are using W3C format also.

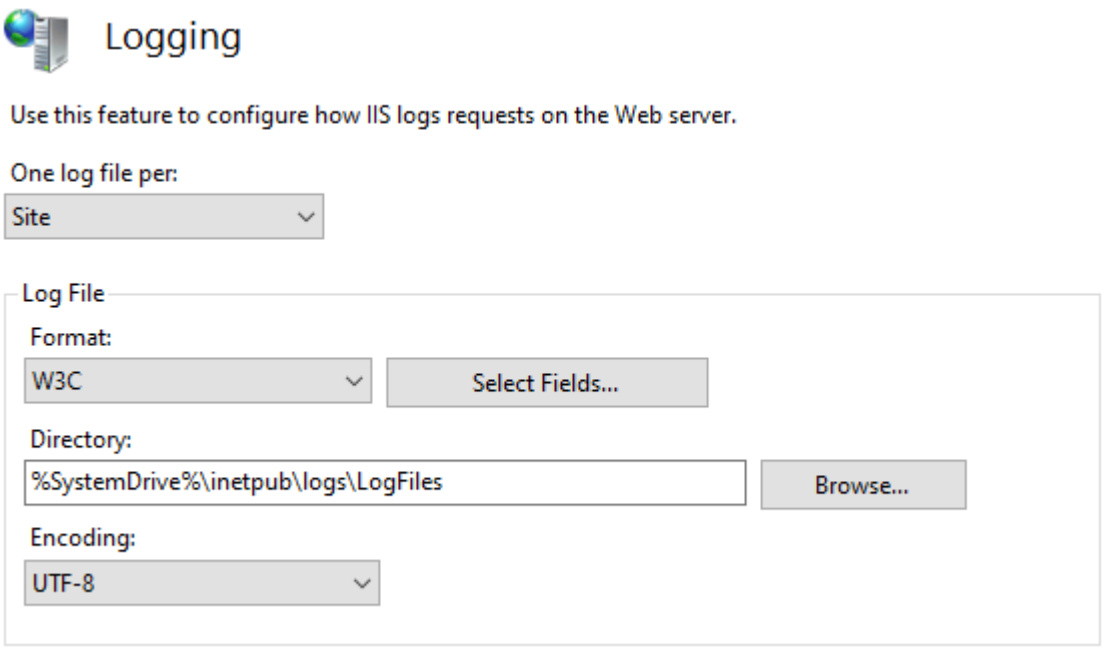


Ensure logs are W3C format

Lastly enable all fields

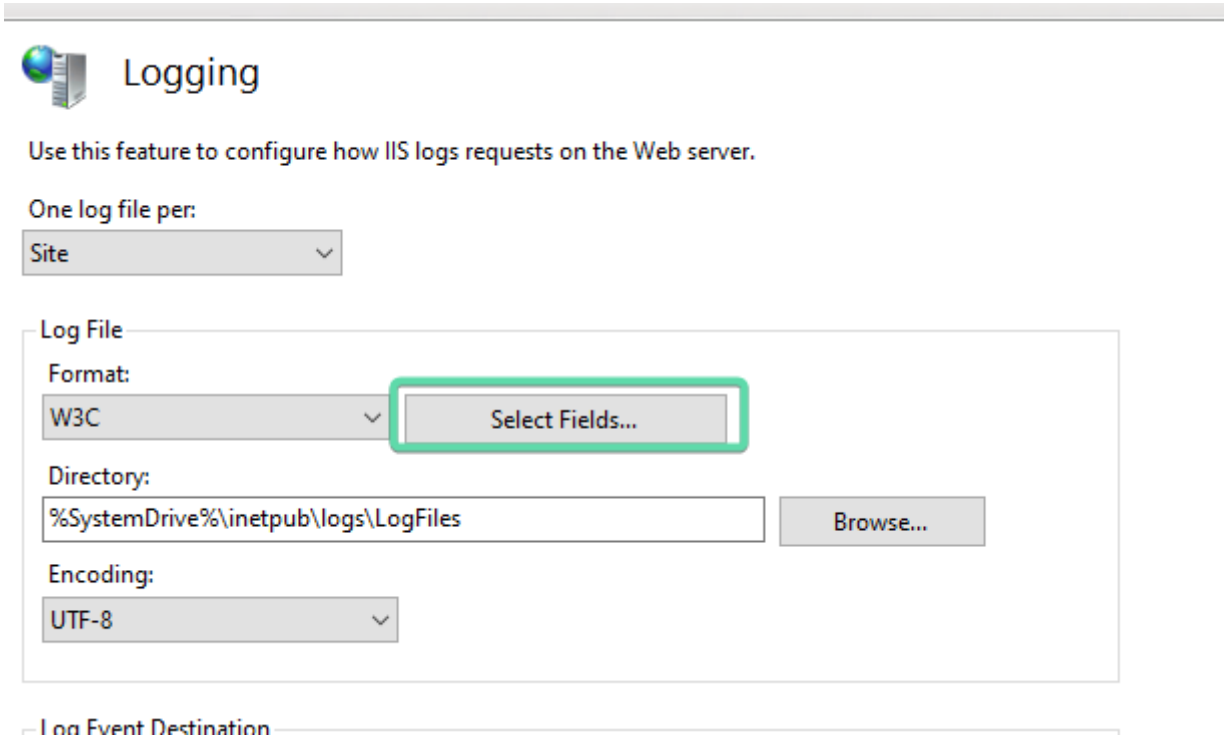


Edit enabled Fields

Go ahead and enable all standard fields

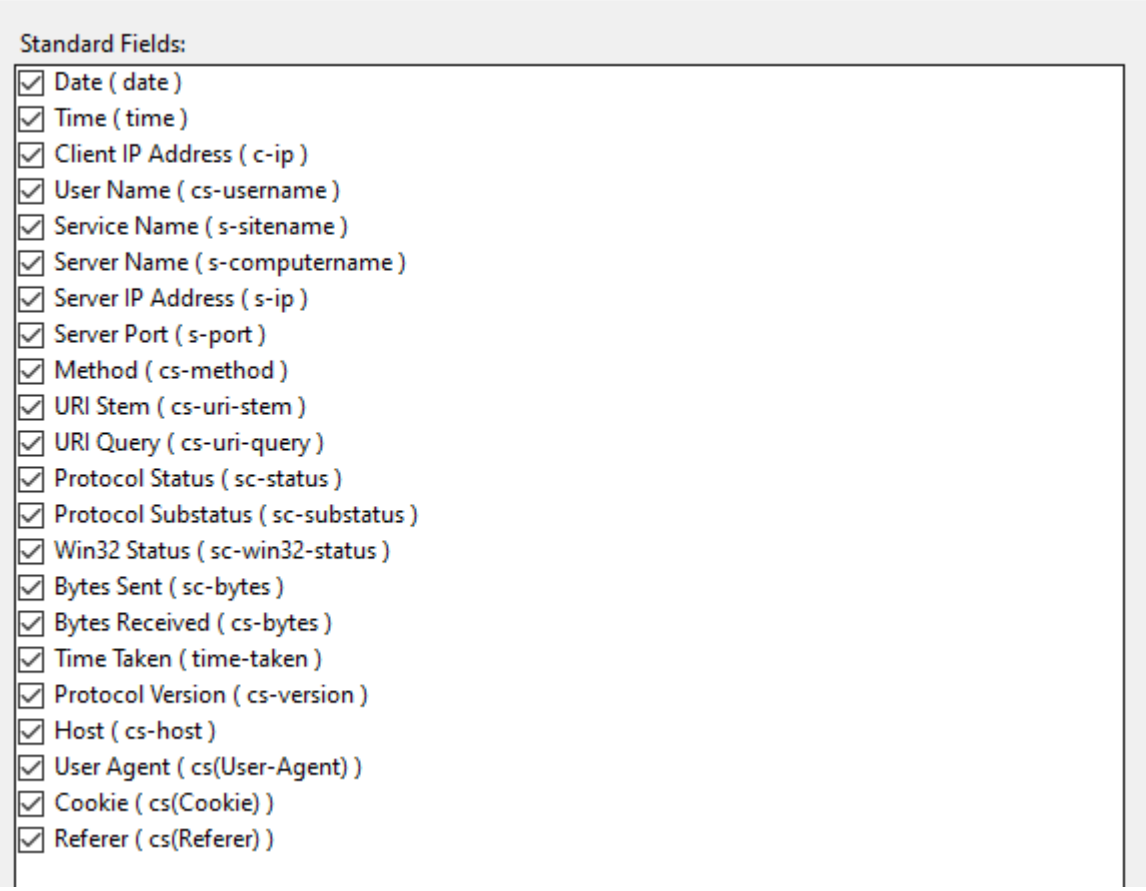


Enable all standard fields

If you have custom fields you will want to take note of them and add them to our GROK pattern latter.

## Setting up the Sidecar Configuration

With our index ready to receive those logs now we need to configure the Graylog Sidecar that will be installed to read our IIS logs.

Using the top menu select **System > Sidecars.**

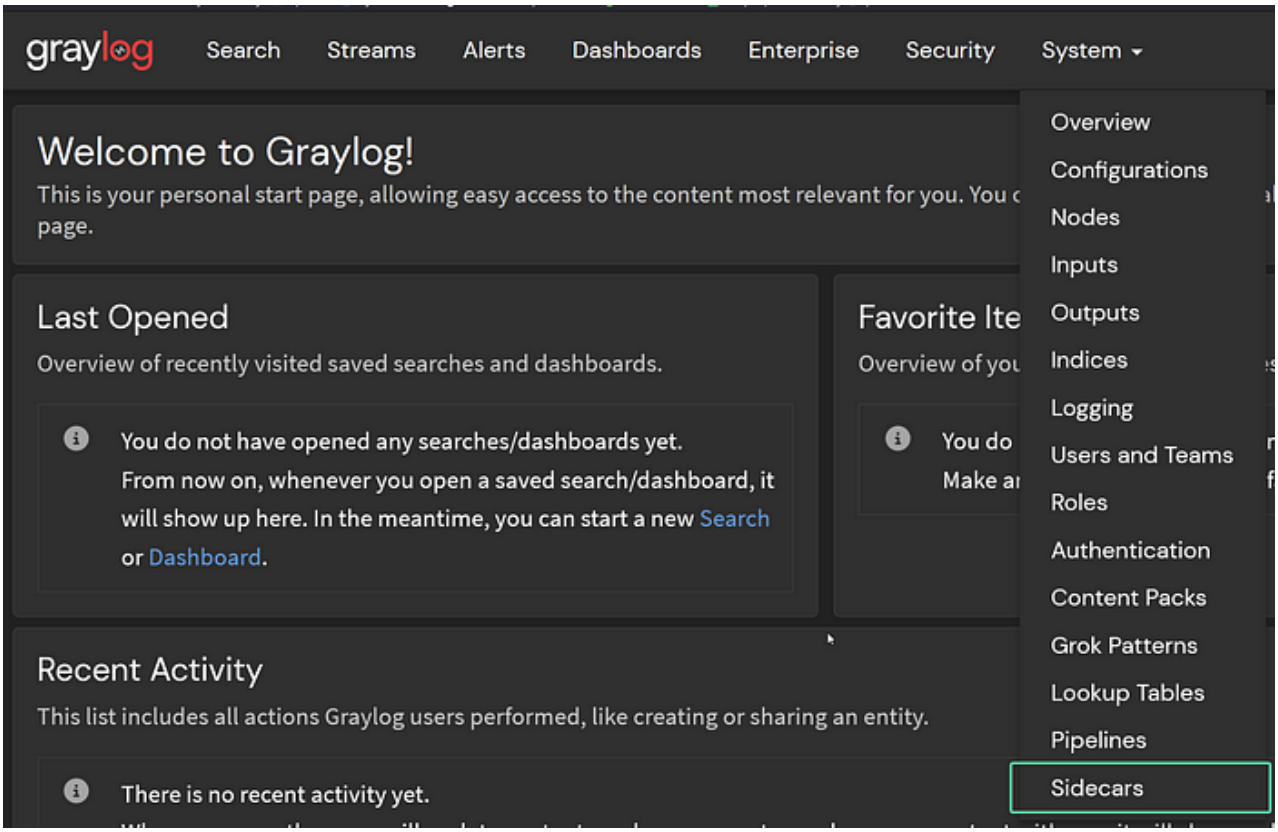


System > Sidecars

Select **configuration** at the from the sidecar menu.

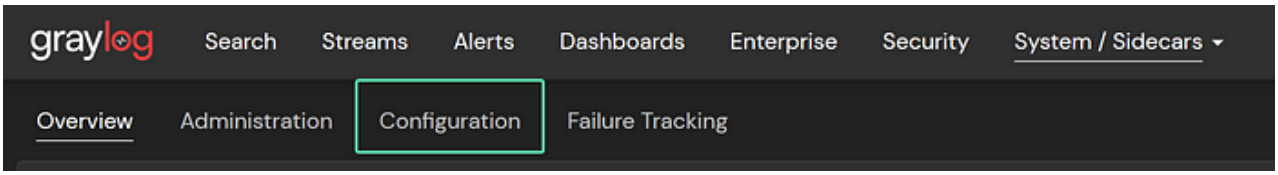


Select Configuration

Create a configuration. We will be creating a **filebeat collector for windows.**

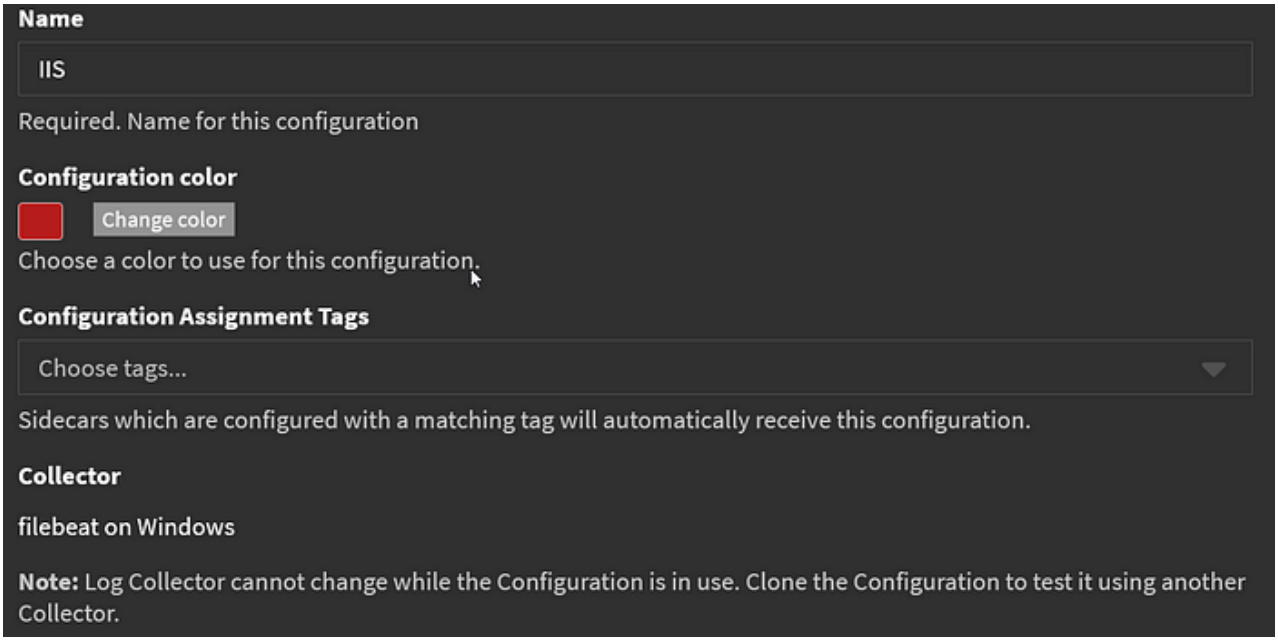


Create the config

The following block of configuration should give the basic outline needed to read any IIS file using the Graylog Sidecar. **Make sure to update the field to fit your logging scheme.**

```
# Needed for Graylog
fields_under_root:true
fields.collector_node_id:${sidecar.nodeName}
fields.gl2_source_collector:${sidecar.nodeId}


output.logstash:
hosts: ["graylog.internal.iso365down.com:5044"]
path:
data:${sidecar.spoolDir!"C:\\ProgramFiles\\Graylog\\sidecar\\cache\\winlogbeat"}\da
logs:${sidecar.spoolDir!"C:\\ProgramFiles\\Graylog\\sidecar"}\logs
```

## Windows Graylog Sidecar Install

Because we are using Graylog 5.2.2 we can use the Graylog Sidecar 1.5.0 code. It can be downloaded **.**

Using the API key we generated in part 4, install the Graylog collector and point it at our Graylog instance.

Make sure you point the system at your Graylog instance using a https://FQDN/api.

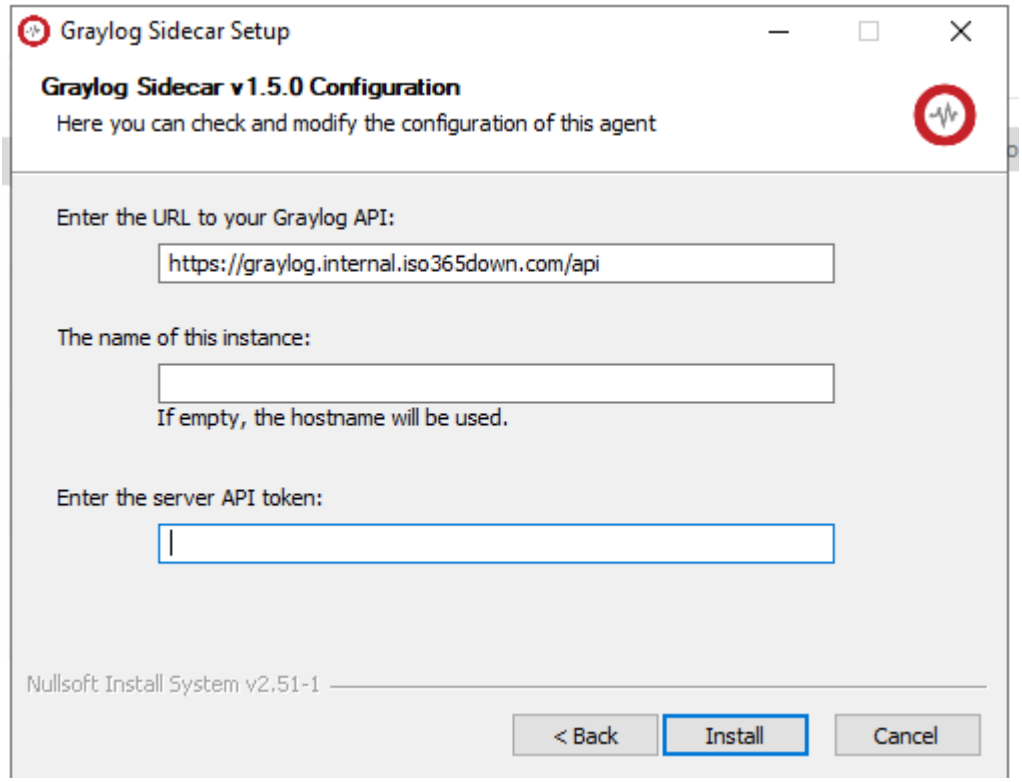


Make sure you use the FQDN of the Graylog Server

And let it install and validate it is showing up in Graylog.

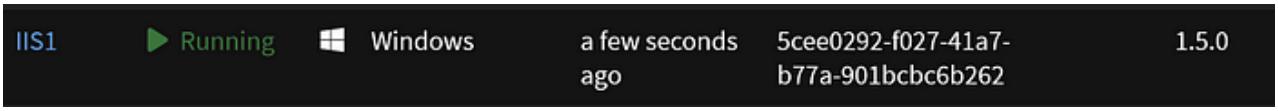


itsworking.jpg

Now go to **Administration** and edit your IIS server and assign your filebeat configuration.

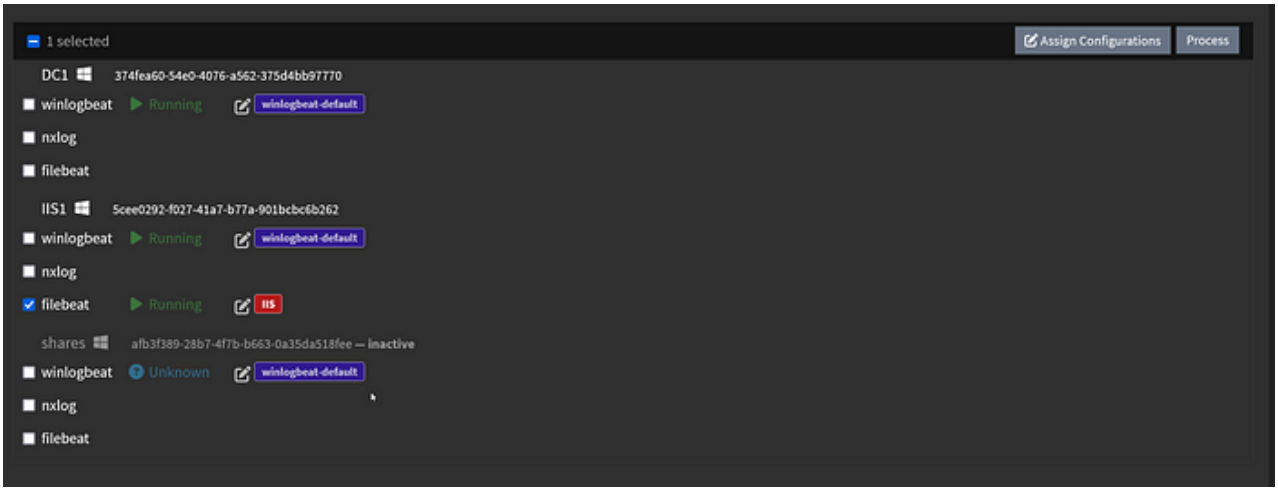


Assign the filebeat configuration

## Creating a new stream

Start by going to **Streams** in the top menu

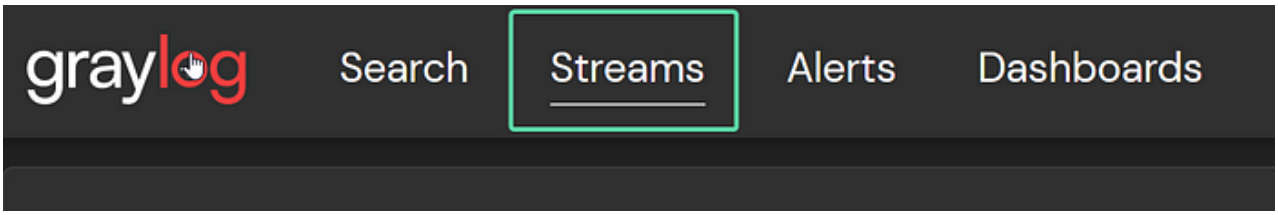


Enter the streams configuration

Create a new stream named IIS Logs and make sure to route the messages to the IIS index set and remove the messages from the default stream.

Name the stream, set the index set it will send data to, and remove the data from the default stream

The newly created stream needs rules configured to send data to it. To do this select **More > Manage Rules.**

Taking the Message ID and Index we saved earlier, load a the message to test against. Search for the field *f*ilebeat_tags and make a rule matching the field name *IIS.*

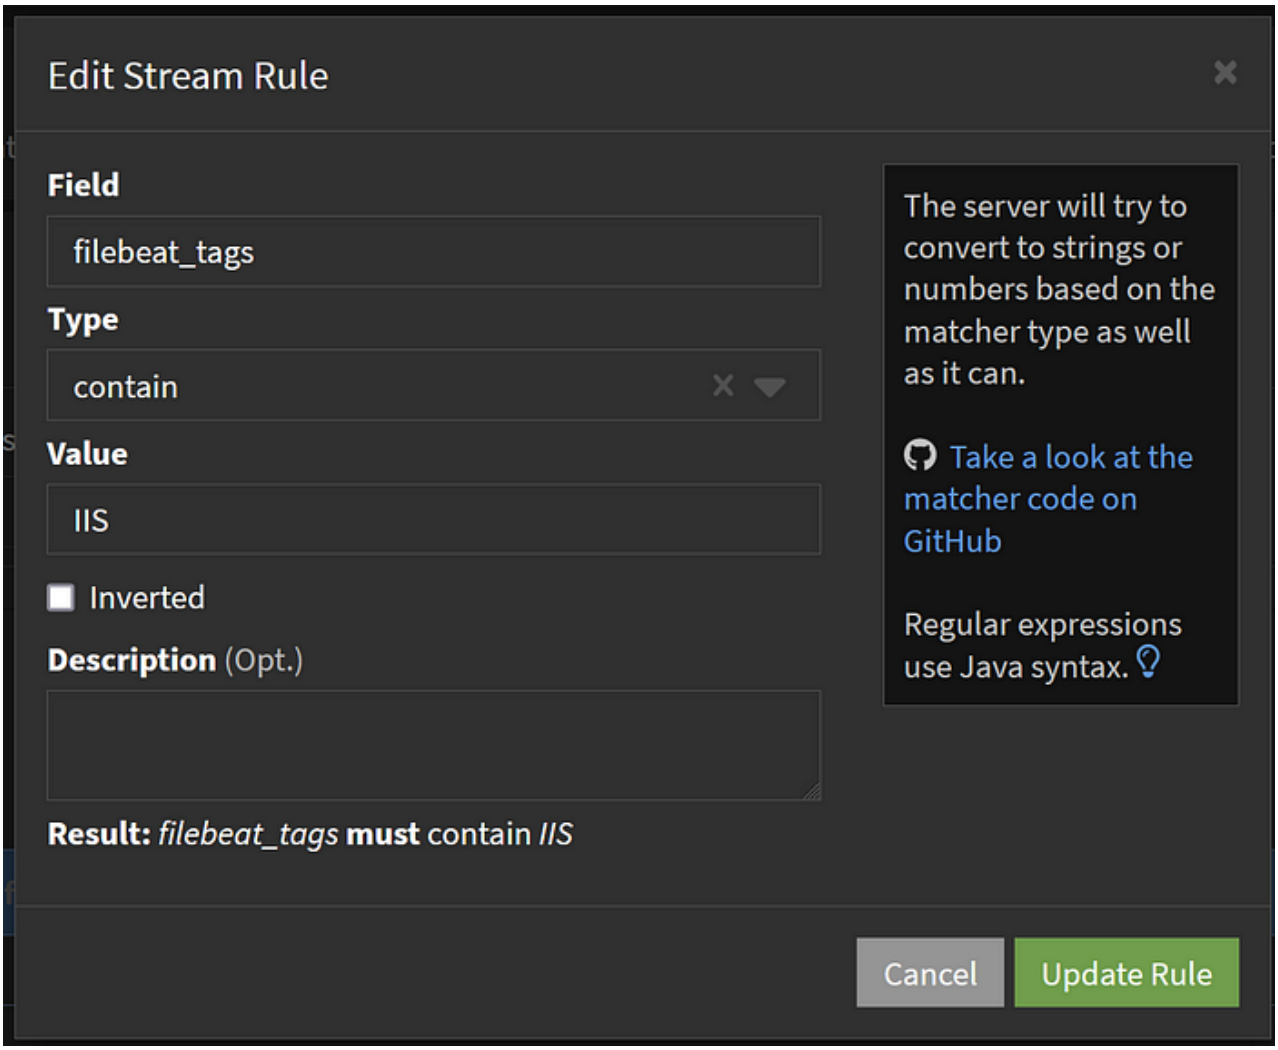


Creating our stream rule

Test the stream rule against the message we selected and make sure the rule matches.

## Decoding our IIS Messages

We have raw IIS logs now, better than what we had before but not super useful. Lets enhance the data we have.

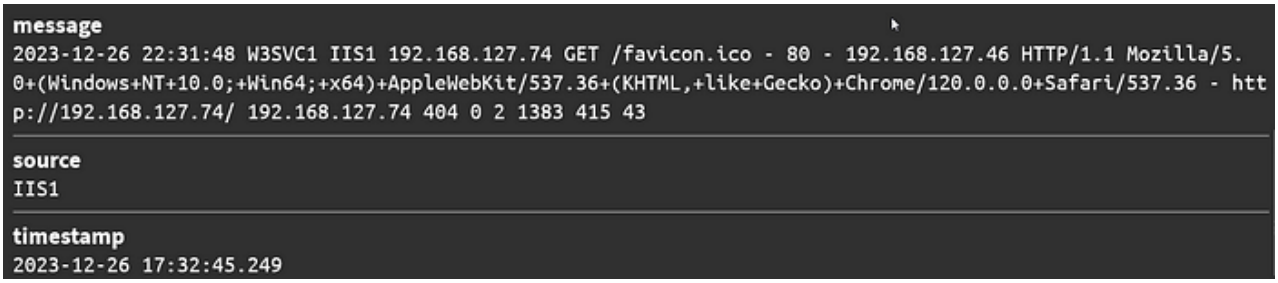


Just a blog of text

Now that we have IIS messages flowing into our IIS Index, gather the message ID and index ID.

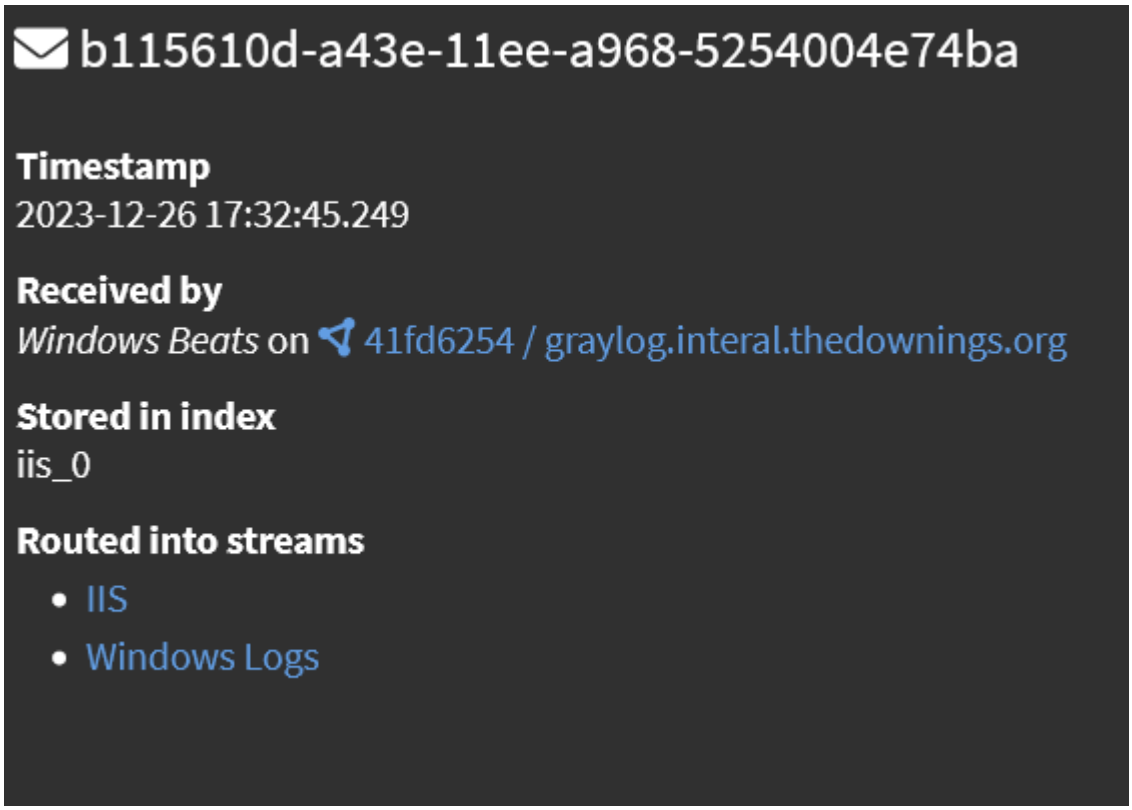


Just to give credit, I am modifying the data found on flatricks github. I found it didn't work anymore and had to make a small modification. (Change %{SERVERNAME:serverName} to %{HOSTNAME:serverName})

To edit our extractor visit **System > Inputs**

Select **Manage Extractors** on our Windows Beat input we created earlier.

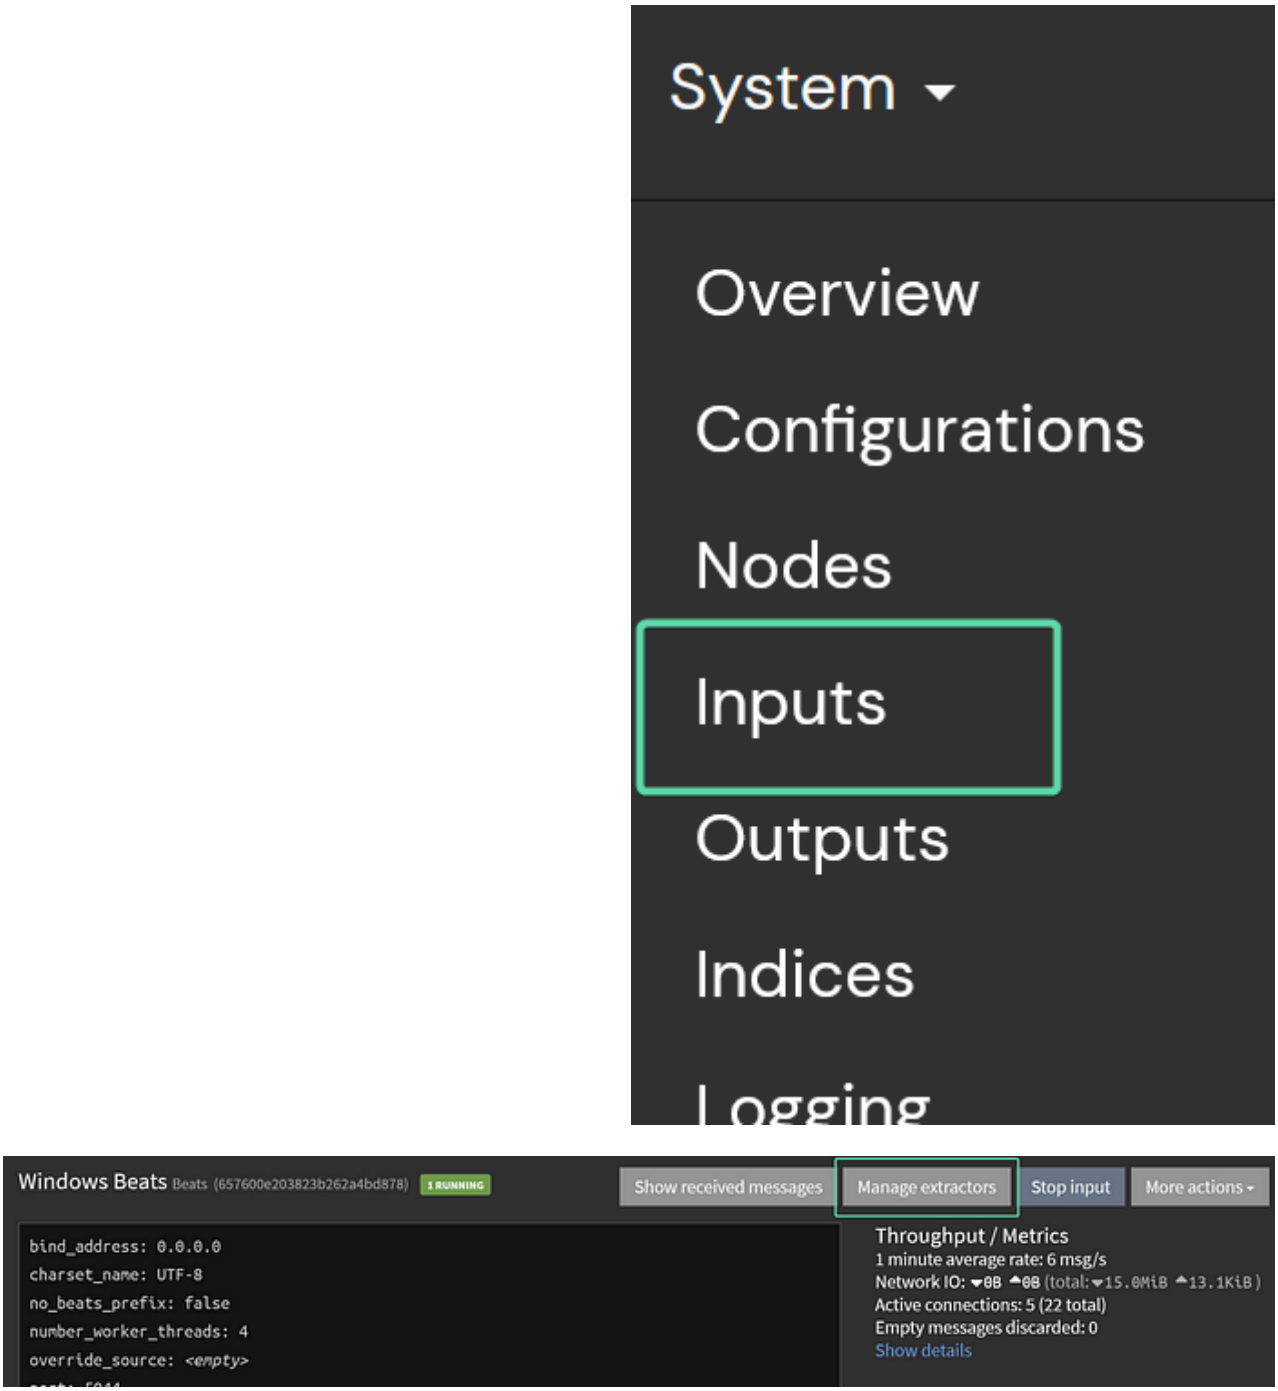


Create a new extractor

I have exported a working extractor.

```
{   [    {       ,       ,       [],       ,       ,       ,       ,       {
,               },       ,           }  ],   }
```

Importing this extractor will create a GUI like this. You may want to adjust the field contains string if your applications uses a different naming scheme.

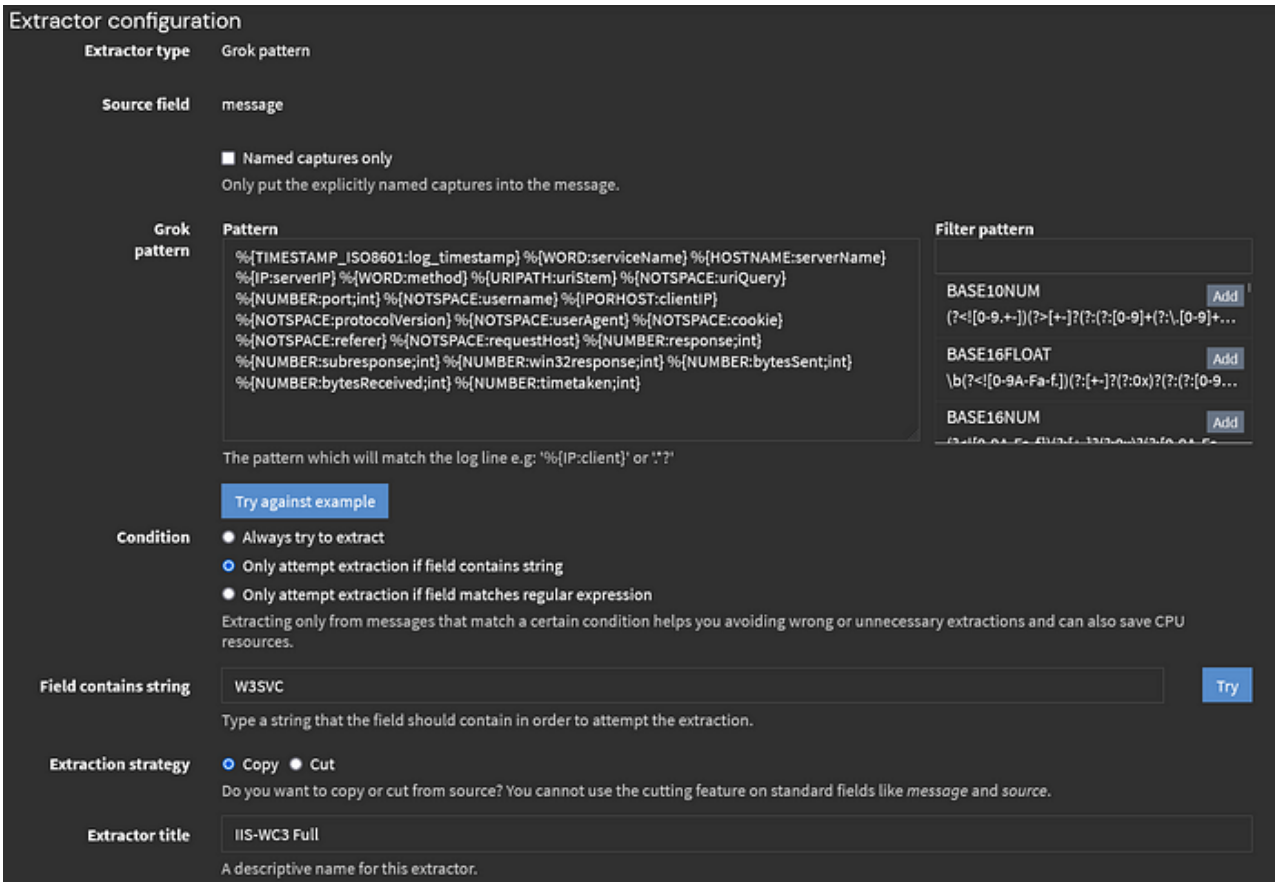


Our import extractor

Here is the raw GROK pattern.

```
%{TIMESTAMP_ISO8601:log_timestamp} %{WORD:serviceName} %{HOSTNAME:serverName} %
{IP:serverIP} %{WORD:method} %{URIPATH:uriStem} %{NOTSPACE:uriQuery} %
{NUMBER:port;int} %{NOTSPACE:username} %{IPORHOST:clientIP} %
{NOTSPACE:protocolVersion} %{NOTSPACE:userAgent} %{NOTSPACE:cookie} %
{NOTSPACE:referer} %{NOTSPACE:requestHost} %{NUMBER:response;int} %
{NUMBER:subresponse;int} %{NUMBER:win32response;int} %{NUMBER:bytesSent;int} %
{NUMBER:bytesReceived;int} %{NUMBER:timetaken;int}
```

If we now generate a new IIS log and test that against our extractor you will see that Graylog has decoded the string of data we sent it.


Successful extraction

If we look at new incoming logs we will see new fields have been added to our message. We can now search on these fields.

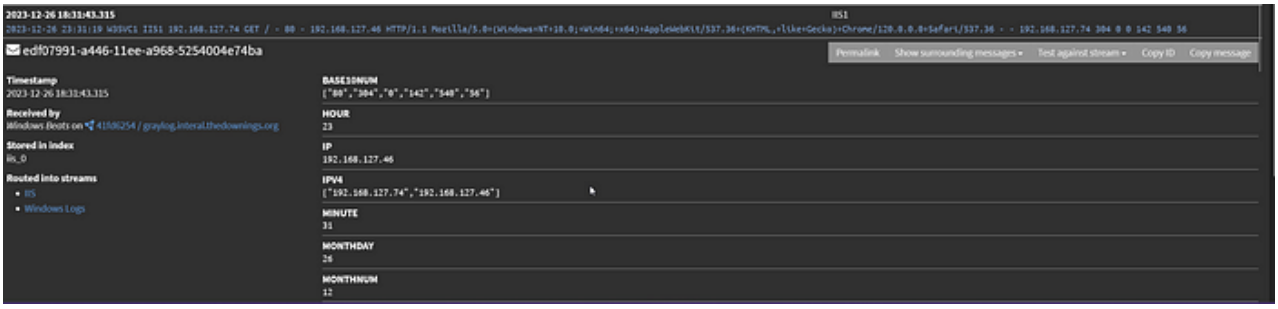


Our enhanced message

Thanks for sticking around! Our next part will be on SQL logs.